

Lab1_FFNN

November 10, 2025

1 Laboratory 1 — Feed Forward Neural Networks (FFNN)

This notebook implements the first lab of the *AI and Cybersecurity* course.

It follows the official brief (`resources/Lab1_FFNN.txt`) and develops a full **Machine Learning pipeline** using PyTorch to explore, train, and evaluate Feed Forward Neural Networks on the CICIDS2017 dataset.

This lab is organized into tasks:

- Task 1: Data preprocessing (cleaning, splitting, outliers, normalization)
- Task 2: Shallow NN (1 layer), train sizes {32, 64, 128}, metrics and analysis; then ReLU change
- Task 3: Impact of specific features (Destination Port), bias test and port removal
- Task 4: Loss function impact (weighted CrossEntropy)
- Task 5: Deep NN, batch size, optimizer comparisons
- Task 6: Overfitting and regularization (dropout, batchnorm, weight decay)

1.1 Setup

```
[1]: import sys
print(sys.executable)
print(sys.version)

/Users/elainnocenti/Documents/Projects/LaTeX/AImSecure/venv/bin/python3.14
3.14.0 (main, Oct 7 2025, 09:34:52) [Clang 17.0.0 (clang-1700.3.19.1)]
```

```
[2]: # --- Import libraries ---
import os
import time
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import json

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScaler, □
    ↵LabelEncoder
from sklearn.utils.class_weight import compute_class_weight
from sklearn.metrics import classification_report, confusion_matrix, □
    ↵accuracy_score, f1_score
```

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader, TensorDataset
```

1.1.1 Device Settings

```
[3]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device = "cpu"
print(f"The device is set to: {device}")
```

The device is set to: cpu

1.1.2 Paths setup

```
[4]: # --- Define Paths ---
project_path = '../'
data_path = project_path + 'data/'
results_path = project_path + 'results/'

# Ensure directories exist
os.makedirs(project_path, exist_ok=True)
os.makedirs(data_path, exist_ok=True)
os.makedirs(results_path, exist_ok=True)

print(f"Project path: {project_path}")
print(f"Data path: {data_path}")
print(f"Results path: {results_path}")
```

Project path: ../
Data path: ../data/
Results path: ../results/

```
[5]: # --- Set visual style ---
sns.set(style="whitegrid", palette="muted", font_scale=1.1)

def save_plot(fig: plt.Figure, filename: str, path: str = "./plots/", fmt: str
             ↪= "png", dpi: int = 300, close_fig: bool = False) -> None:
    """
    Save a Matplotlib figure in a specific directory.

    Args:
        fig (plt.Figure): Matplotlib figure object to save.
        filename (str): Name of the file to save (e.g., 'plot.png').
        path (str, optional): Directory path to save the figure. Defaults to './
                             ↪plots/'.
        fmt (str, optional): File format for the saved figure. Defaults to
                            ↪'png'.
```

```

    dpi (int, optional): Dots per inch for the saved figure. Defaults to
    ↪300.

    Returns:
        None
    """
    # Ensure the directory exists
    os.makedirs(path, exist_ok=True)
    save_path = os.path.join(path, f"{filename}.{fmt}")

    # Save the figure
    fig.savefig(save_path, bbox_inches='tight', pad_inches=0.1, dpi=dpi, u
    ↪format=fmt)
    # plt.close(fig) # Removed to display plots in notebook

    if close_fig:
        plt.close(fig)

    print(f"Saved plot: {save_path}")

```

1.2 Task 1 — Data preprocessing

In this task, we prepare the dataset and ensure proper data quality and normalization before training.

We will:

- Load the dataset from `lab/data/dataset_lab_1.csv`.
- Inspect basic statistics, feature ranges, and class distribution.
- Remove missing values and duplicates (report before/after counts).
- Split the data into **train / validation / test (60% / 20% / 20%)**, using a fixed random seed for reproducibility.
- Perform **outlier analysis** (e.g., KDE, Z-score, IQR) and decide on an appropriate **scaling method**.
- Fit the scaler on the training data only and apply it to validation and test splits.

```
[6]: # --- Load dataset and perform initial inspection ---

# Set random seed for reproducibility
np.random.seed(42)

# Create directory for plots
save_dir = results_path + 'images/' + 'task1_plots/'
os.makedirs(save_dir, exist_ok=True)

# Load Dataset
file_path = data_path + 'dataset_lab_1.csv'
df = pd.read_csv(file_path)

# Basic info
print("Shape (raw):", df.shape)
print("Columns:", list(df.columns))
```

```
print("\nLabel distribution (raw):")
print(df['Label'].value_counts(dropna=False))
```

Shape (raw): (31507, 17)
Columns: ['Flow Duration', 'Flow IAT Mean', 'Fwd PSH Flags', 'Bwd Packet Length Mean', 'Bwd Packet Length Max', 'Flow Bytes/s', 'Down/Up Ratio', 'SYN Flag Count', 'Fwd Packet Length Mean', 'Fwd IAT Std', 'Packet Length Mean', 'Fwd Packet Length Max', 'Subflow Fwd Packets', 'Flow Packets/s', 'Total Fwd Packets', 'Destination Port', 'Label']

Label distribution (raw):

Label
Benign 20000
DoS Hulk 5000
PortScan 5000
Brute Force 1507
Name: count, dtype: int64

[7]: df

```
[7]:      Flow Duration  Flow IAT Mean  Fwd PSH Flags  Bwd Packet Length Mean \
0          303376     3.033760e+04           0             749.4
1            117     1.170000e+02           0              0.0
2            142     1.420000e+02           0              0.0
3            191     6.366667e+01           0             52.0
4              4     4.000000e+00           0              0.0
...
31502      5710955    1.903652e+06           0              0.0
31503      5862561    1.954187e+06           0              0.0
31504      5854872    1.951624e+06           0              0.0
31505      5555547    1.851849e+06           0              0.0
31506      5018204    1.672735e+06           0              0.0

      Bwd Packet Length Max  Flow Bytes/s  Down/Up Ratio  SYN Flag Count \
0            1448        12743.263           0              0
1              0         0.000             1              0
2              0         84507.040           0              0
3              52         921465.940           1              0
4              0        3000000.000           0              0
...
31502          ...          ...           ...           ...
31503          ...          ...           ...           ...
31504          ...          ...           ...           ...
31505          ...          ...           ...           ...
31506          ...          ...           ...           ...

      Fwd Packet Length Mean  Fwd IAT Std  Packet Length Mean \

```

0	19.833334	98776.15	322.16666
1	0.000000	0.00	0.00000
2	6.000000	0.00	6.00000
3	36.000000	0.00	42.40000
4	6.000000	0.00	6.00000
...
31502	0.000000	4037277.80	0.00000
31503	0.000000	4144374.80	0.00000
31504	0.000000	4139029.80	0.00000
31505	0.000000	3927356.50	0.00000
31506	0.000000	3547404.80	0.00000

	Fwd Packet Length	Max	Subflow Fwd Packets	Flow Packets/s	\
0	119		6	36.258636	
1	0		1	17094.018000	
2	6		2	14084.507000	
3	36		2	20942.408000	
4	6		2	500000.000000	
...
31502	0		3	0.700408	
31503	0		3	0.682296	
31504	0		3	0.683192	
31505	0		3	0.720001	
31506	0		3	0.797098	

	Total Fwd Packets	Destination Port	Label
0	6	443	Benign
1	1	52631	Benign
2	2	80	Benign
3	2	53	Benign
4	2	49467	Benign
...
31502	3	80	Brute Force
31503	3	80	Brute Force
31504	3	80	Brute Force
31505	3	80	Brute Force
31506	3	80	Brute Force

[31507 rows x 17 columns]

```
[8]: def plot_class_distribution(
    df,
    name_fig='class_distribution',
    label_col='Label',
    save_path='./plots/',
    fig_size=(8, 5),
    palette='pastel'
```

```

):
"""
Plot the number of samples for each class.

Args:
    df (pd.DataFrame): The input DataFrame.
    label_col (str): The name of the label column. Defaults to 'Label'.
    save_path (str): The directory to save the plot. Defaults to './plots/'.
    fig_size (tuple): Size of the figure. Defaults to (8, 5).
    palette (str): Seaborn palette.

os.makedirs(save_path, exist_ok=True)

# Create a figure and axes for the plot
fig, ax = plt.subplots(figsize=fig_size)

# Create a countplot showing the distribution of classes
sns.countplot(x=label_col, data=df, order=df[label_col].value_counts() .
index, ax=ax, palette=palette, hue=label_col, legend=False)

# Set the title and labels for the plot
ax.set_title("Class Distribution")
ax.set_xlabel("Traffic Type")
ax.set_ylabel("Count")

# Rotate x-axis labels for better readability
plt.xticks(rotation=30)

# Annotate bars with counts
for p in ax.patches:
    height = p.get_height()
    ax.annotate(f'{height}', (p.get_x() + p.get_width() / 2., height),
                ha='center', va='bottom', fontsize=10)

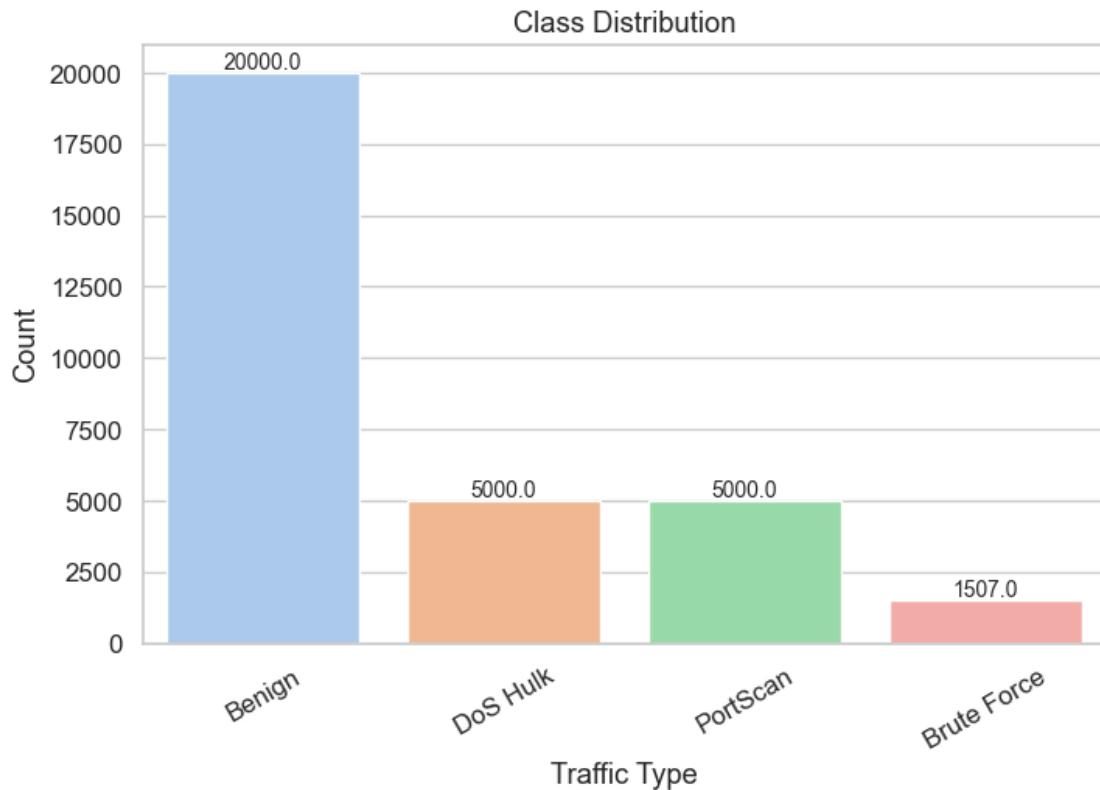
# Save the plot to the specified path
save_plot(fig, name_fig, save_path)

```

[9]: # --- Visualize raw data characteristics and distributions ---

```
# Plot class distribution to understand data balance
plot_class_distribution(df, 'class_distribution', save_path=save_dir)
```

Saved plot: ../results/images/task1_plots/class_distribution.png



1.2.1 Removing NaN and duplicates

```
[10]: # --- Remove NaN and duplicates ---

raw_n = len(df)

# Drop NaN
df = df.dropna()
clean_n = len(df)
print(f"Removed {raw_n-clean_n} rows (NaN)")

# Drop duplicates
df = df.drop_duplicates()
clean_n = len(df)
print(f"Removed {raw_n-clean_n} rows (duplicates)")

print(f"New shape: {df.shape}")
```

Removed 20 rows (NaN)
 Removed 2114 rows (duplicates)
 New shape: (29393, 17)

```
[11]: # --- Handle infinite values ---

# Replace infinite values with NaN
df.replace([np.inf, -np.inf], np.nan, inplace=True)

# Check for and report the number of infinite values (now NaNs)
inf_counts = df.isnull().sum()
print("\nNumber of infinite values (replaced with NaN) per column:")
print(inf_counts[inf_counts > 0])

# Drop rows with NaN values (including those that were originally infinite)
initial_rows = len(df)
df.dropna(inplace=True)
rows_after_inf_nan_drop = len(df)

print(f"\nRemoved {initial_rows - rows_after_inf_nan_drop} rows containing"
      " infinite or NaN values. New shape: {df.shape}")
```

Number of infinite values (replaced with NaN) per column:
Flow Bytes/s 7
Flow Packets/s 7
dtype: int64

Removed 7 rows containing infinite or NaN values. New shape: (29386, 17)

Q: How many samples did you have before and after removing missing and duplicates entries? We initially had **31507** samples bold text, and after removing missing, duplicate, and infinite entries, we obtained 29,386 clean samples. In total, **2121** (2114 missing and duplicates + 7 infinite values) rows **were removed** during the data cleaning process.

```
[12]: print("\nLabel distribution (after NaN and duplicates removal):")
print(df['Label'].value_counts())
```

Label distribution (after NaN and duplicates removal):
Label
Benign 19242
PortScan 4849
DoS Hulk 3868
Brute Force 1427
Name: count, dtype: int64

```
[13]: # Encode labels
label_encoder = LabelEncoder()
df['Label'] = label_encoder.fit_transform(df['Label'])

df.head(10)
```

[13]:	Flow Duration	Flow IAT Mean	Fwd PSH Flags	Bwd Packet Length Mean	\
0	303376	3.033760e+04	0	749.400000	
1	117	1.170000e+02	0	0.000000	
2	142	1.420000e+02	0	0.000000	
3	191	6.366667e+01	0	52.000000	
4	4	4.000000e+00	0	0.000000	
5	162322	3.959073e+03	0	19.416666	
6	61050653	2.907174e+06	0	22.818182	
7	108	1.080000e+02	0	0.000000	
8	764920	7.649200e+05	0	0.000000	
9	122255	6.112750e+04	0	0.000000	
	Bwd Packet Length Max	Flow Bytes/s	Down/Up Ratio	SYN Flag Count	\
0	1448	1.274326e+04	0	0	
1	0	0.000000e+00	1	0	
2	0	8.450704e+04	0	0	
3	52	9.214659e+05	1	0	
4	0	3.000000e+06	0	0	
5	51	4.102956e+03	1	0	
6	233	5.230083e+01	1	0	
7	0	1.111111e+05	0	0	
8	0	1.568792e+01	0	0	
9	0	0.000000e+00	0	0	
	Fwd Packet Length Mean	Fwd IAT Std	Packet Length Mean	\	
0	19.833334	98776.150	322.166660		
1	0.000000	0.000	0.000000		
2	6.000000	0.000	6.000000		
3	36.000000	0.000	42.400000		
4	6.000000	0.000	6.000000		
5	11.111111	30630.129	15.488372		
6	267.454560	5145240.000	138.826080		
7	6.000000	0.000	6.000000		
8	6.000000	0.000	6.000000		
9	0.000000	0.000	0.000000		
	Fwd Packet Length Max	Subflow Fwd Packets	Flow Packets/s	\	
0	119	6	36.258636		
1	0	1	17094.018000		
2	6	2	14084.507000		
3	36	2	20942.408000		
4	6	2	500000.000000		
5	43	18	258.744960		
6	1460	11	0.360357		
7	6	2	18518.518000		
8	6	2	2.614652		
9	0	2	24.538874		

	Total Fwd Packets	Destination Port	Label
0	6	443	0
1	1	52631	0
2	2	80	0
3	2	53	0
4	2	49467	0
5	18	21	0
6	11	80	0
7	2	80	0
8	2	443	0
9	2	443	0

```
[14]: # --- Display the label encoding mapping ---
print("\nLabel Encoding Mapping:")
for i, label in enumerate(label_encoder.classes_):
    print(f"[label]: {i}")
```

Label Encoding Mapping:

Benign: 0
Brute Force: 1
DoS Hulk: 2
PortScan: 3

```
[15]: # Checking again null values
print(df.isnull().sum())
```

Flow Duration	0
Flow IAT Mean	0
Fwd PSH Flags	0
Bwd Packet Length Mean	0
Bwd Packet Length Max	0
Flow Bytes/s	0
Down/Up Ratio	0
SYN Flag Count	0
Fwd Packet Length Mean	0
Fwd IAT Std	0
Packet Length Mean	0
Fwd Packet Length Max	0
Subflow Fwd Packets	0
Flow Packets/s	0
Total Fwd Packets	0
Destination Port	0
Label	0

dtype: int64

```
[16]: # Checking again for the infinite values
print(df.isin([np.inf, -np.inf]).sum())
```

```
Flow Duration      0
Flow IAT Mean     0
Fwd PSH Flags     0
Bwd Packet Length Mean 0
Bwd Packet Length Max 0
Flow Bytes/s      0
Down/Up Ratio     0
SYN Flag Count    0
Fwd Packet Length Mean 0
Fwd IAT Std       0
Packet Length Mean 0
Fwd Packet Length Max 0
Subflow Fwd Packets 0
Flow Packets/s    0
Total Fwd Packets 0
Destination Port   0
Label              0
dtype: int64
```

1.2.2 Split data

We use a **stratified approach** due to the class imbalance.

```
[17]: # --- Split data ---

# Split features/target
label_col = 'Label'
feature_cols = [c for c in df.columns if c != label_col]
X = df[feature_cols].values
y = df[label_col].values

# Train/val/test split 60/20/20 with stratify
X_train, X_tmp, y_train, y_tmp = train_test_split(
    X, y, test_size=0.4, stratify=y, random_state=42
)
X_val, X_test, y_val, y_test = train_test_split(
    X_tmp, y_tmp, test_size=0.5, stratify=y_tmp, random_state=42
)

print("Data Splits:")
print(f"Train set: {X_train.shape[0]} samples")
print(f"Validation set: {X_val.shape[0]} samples")
print(f"Test set: {X_test.shape[0]} samples\n")

def print_label_counts(name, y):
```

```

labels, counts = np.unique(y, return_counts=True)
count_width = 6
print(f"{'name':<8}", end=" ")
for label, count in zip(labels, counts):
    print(f"{'label':>{count_width},}", end="   ")
print()

print_label_counts("Train", y_train)
print_label_counts("Val", y_val)
print_label_counts("Test", y_test)

```

Data Splits:

Train set: 17,631 samples
Validation set: 5,877 samples
Test set: 5,878 samples

Train	0: 11,545	1: 856	2: 2,321	3: 2,909
Val	0: 3,848	1: 286	2: 773	3: 970
Test	0: 3,849	1: 285	2: 774	3: 970

1.2.3 Outliers Detection

```
[18]: # --- Outlier detection (Z-score and IQR) ---

# Convert X_train to DataFrame to use select_dtypes
X_train_df = pd.DataFrame(X_train, columns=feature_cols)

num_cols = X_train_df.select_dtypes(include=[np.number]).columns.tolist()
print(num_cols)

# Z-score method
zs = ((X_train_df[num_cols] - X_train_df[num_cols].mean()) / 
      X_train_df[num_cols].std(ddof=0)).abs()
outlier_counts_z = (zs > 3).sum().sort_values(ascending=False)
print("\n[Z-SCORE OUTLIERS] features:")
print(outlier_counts_z)

# IQR method
Q1 = X_train_df[num_cols].quantile(0.25)
Q3 = X_train_df[num_cols].quantile(0.75)
IQR = Q3 - Q1
outliers_iqr = ((X_train_df[num_cols] < (Q1 - 1.5 * IQR)) | 
                  (X_train_df[num_cols] > (Q3 + 1.5 * IQR))).sum()
outlier_counts_iqr = outliers_iqr.sort_values(ascending=False)
print("\n[IQR OUTLIERS] features:")
print(outlier_counts_iqr)
```

['Flow Duration', 'Flow IAT Mean', 'Fwd PSH Flags', 'Bwd Packet Length Mean',

```
'Bwd Packet Length Max', 'Flow Bytes/s', 'Down/Up Ratio', 'SYN Flag Count', 'Fwd Packet Length Mean', 'Fwd IAT Std', 'Packet Length Mean', 'Fwd Packet Length Max', 'Subflow Fwd Packets', 'Flow Packets/s', 'Total Fwd Packets', 'Destination Port']
```

[Z-SCORE OUTLIERS] features:

```
Fwd PSH Flags      708  
SYN Flag Count    708  
Fwd IAT Std       567  
Bwd Packet Length Max 523  
Flow Packets/s    522  
Bwd Packet Length Mean 340  
Packet Length Mean 272  
Destination Port   247  
Fwd Packet Length Max 238  
Flow IAT Mean      207  
Fwd Packet Length Mean 157  
Subflow Fwd Packets 88  
Total Fwd Packets  88  
Down/Up Ratio      48  
Flow Bytes/s       37  
Flow Duration       0  
dtype: int64
```

[IQR OUTLIERS] features:

```
Bwd Packet Length Max 3951  
Destination Port      3879  
Flow Duration         3443  
Bwd Packet Length Mean 3180  
Packet Length Mean    2982  
Flow IAT Mean          2913  
Flow Bytes/s           2826  
Fwd IAT Std            2662  
Subflow Fwd Packets   1467  
Total Fwd Packets     1467  
Flow Packets/s         1417  
Fwd Packet Length Mean 997  
Fwd Packet Length Max 929  
Fwd PSH Flags          708  
SYN Flag Count          708  
Down/Up Ratio            48  
dtype: int64
```

```
[19]: # --- Visualize raw distributions for selected features ---
```

```
selected_features = [  
    'Flow Duration', 'Flow IAT Mean', 'Bwd Packet Length Mean',
```

```

'Bwd Packet Length Max', 'Flow Bytes/s', 'Fwd Packet Length Mean',
'Fwd IAT Std', 'Packet Length Mean', 'Fwd Packet Length Max', 'Flow Packets/
↳S'
]

# Convert X_train to DataFrame for easier plotting with column names
X_train_df = pd.DataFrame(X_train, columns=feature_cols)

def plot_distributions(df, features, title_prefix, save_name):
    n = len(features)
    fig, axes = plt.subplots(nrows=n, ncols=2, figsize=(10, 3 * n))

    for i, col in enumerate(features):
        # KDE Plot
        sns.kdeplot(df[col], ax=axes[i, 0], fill=True, color="steelblue")
        axes[i, 0].set_title(f"{title_prefix} - KDE: {col}")
        axes[i, 0].set_xlabel("")

        # Boxplot
        sns.boxplot(x=df[col], ax=axes[i, 1], color="lightcoral")
        axes[i, 1].set_title(f"{title_prefix} - Boxplot: {col}")
        axes[i, 1].set_xlabel("")

    plt.tight_layout()

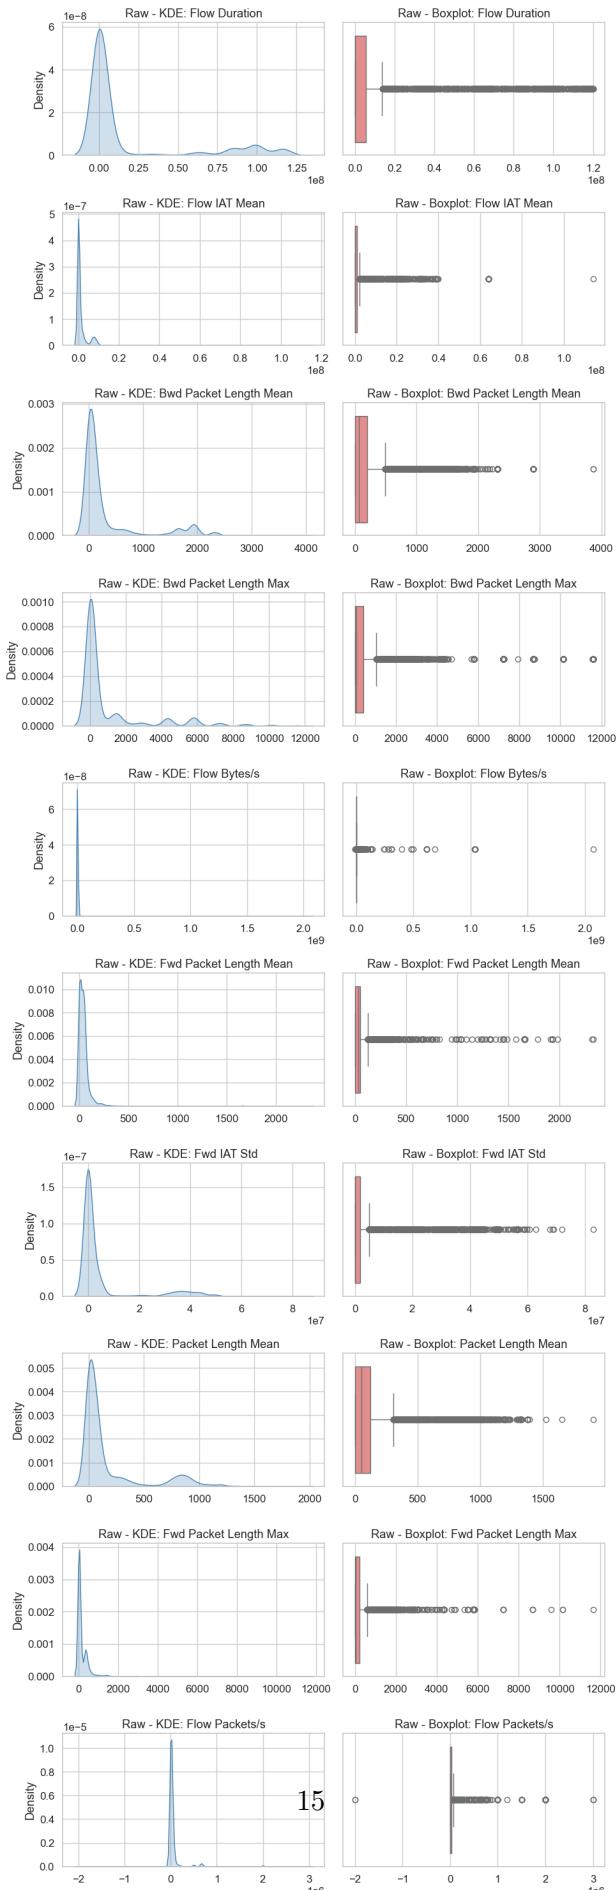
# Save the plot to the specified path
save_plot(fig, save_name, save_dir, fmt='pdf')

plt.show()

plot_distributions(X_train_df, selected_features, "Raw", "raw_distributions")

```

Saved plot: ../results/images/task1_plots/raw_distributions.pdf



1.2.4 Data Normalization

```
[20]: # Standardize the features (StandardScaler)
scaler1 = StandardScaler()
X_train_std = scaler1.fit_transform(X_train)
X_val_std = scaler1.transform(X_val)
X_test_std = scaler1.transform(X_test)

# Standardize the features (RobustScaler)
scaler2 = RobustScaler()
X_train_rob = scaler2.fit_transform(X_train)
X_val_rob = scaler2.transform(X_val)
X_test_rob = scaler2.transform(X_test)
```

```
[21]: # --- Visual comparison of Standard vs Robust scaling ---

# Convert scaled arrays back to DataFrames for easier plotting with column names
X_train_std_df = pd.DataFrame(X_train_std, columns=feature_cols)
X_train_rob_df = pd.DataFrame(X_train_rob, columns=feature_cols)

# Adjust subplot grid to accommodate all selected features (10 features, 5 rows
# x 2 columns)
fig, axes = plt.subplots(5, 2, figsize=(10, 20)) # Increased figsize for better
# readability

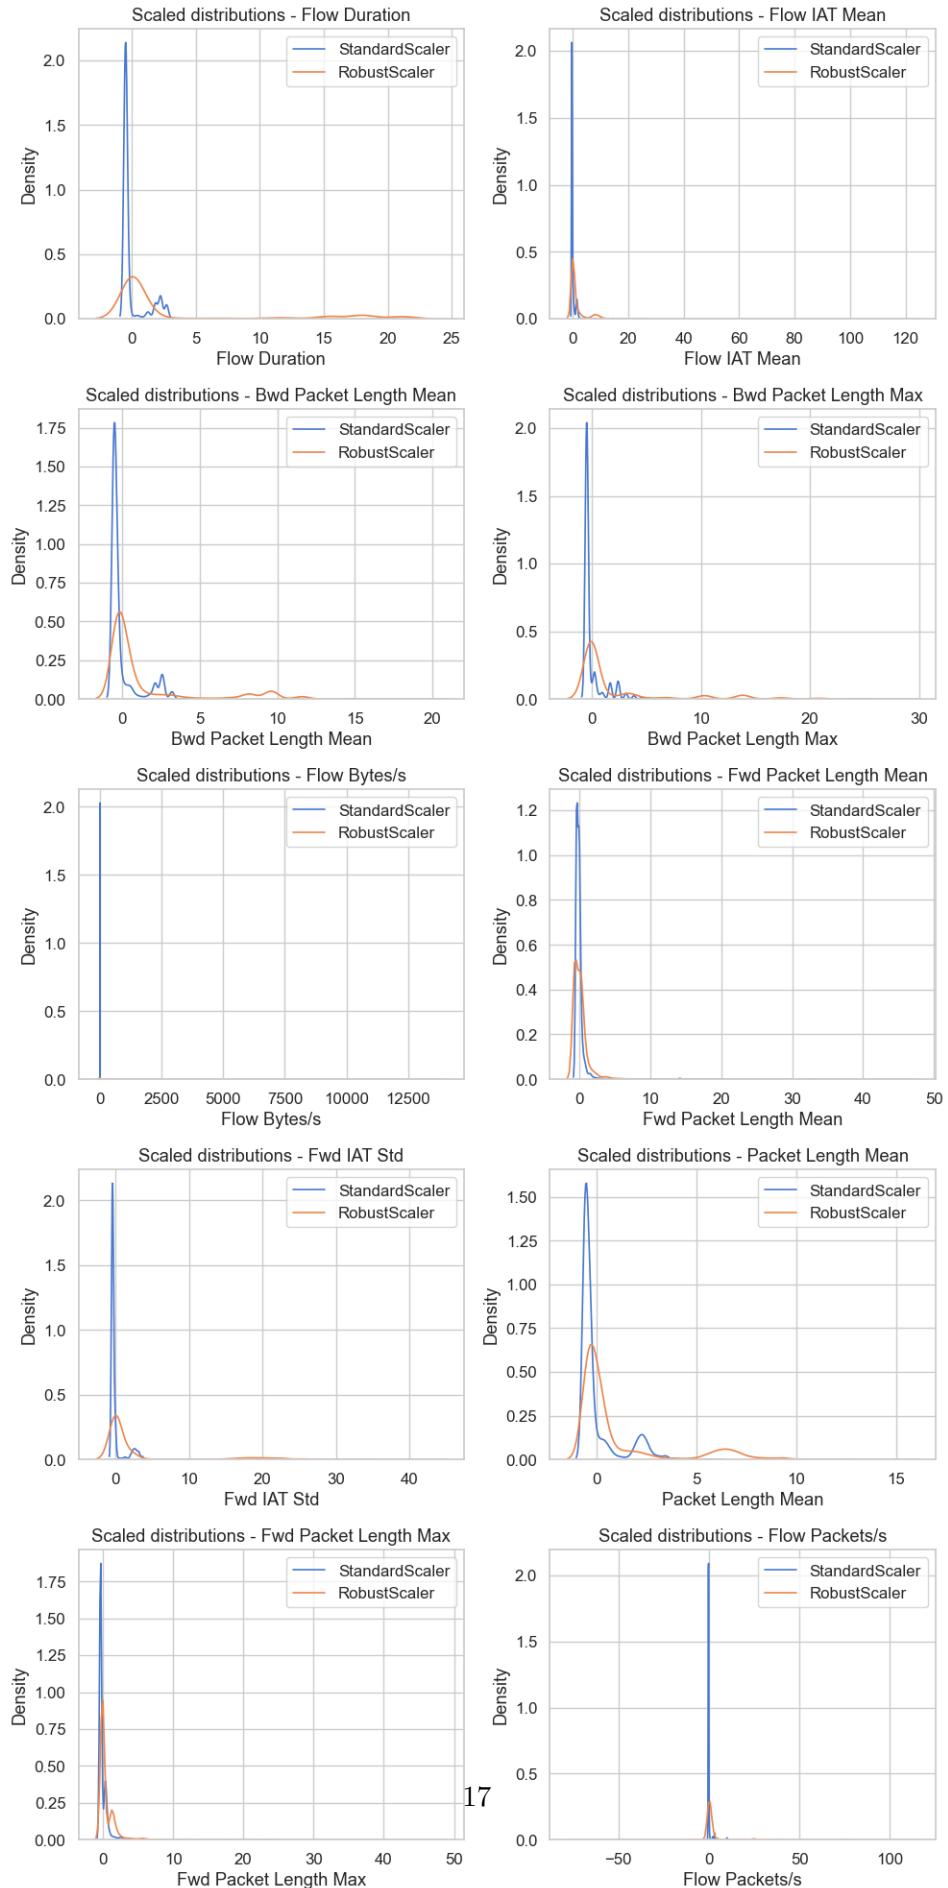
# Plot all selected features for comparison
for i, col in enumerate(selected_features):
    row_idx = i // 2
    col_idx = i % 2
    sns.kdeplot(X_train_std_df[col], ax=axes[row_idx, col_idx], color="blue",
                 label="StandardScaler", lw=1.2)
    sns.kdeplot(X_train_rob_df[col], ax=axes[row_idx, col_idx], color="red",
                 label="RobustScaler", lw=1.2)
    axes[row_idx, col_idx].set_title(f"Scaled distributions - {col}")
    axes[row_idx, col_idx].legend()

plt.tight_layout()

# Save the plot to the specified path
save_plot(fig, 'raw_distributions_comparison', save_dir, fmt='pdf')

plt.show()
plt.close(fig)
```

Saved plot: ../results/images/task1_plots/raw_distributions_comparison.pdf



```
[22]: # --- Compare scaling statistics numerically ---

stats_std = X_train_std_df.describe().T[['mean', 'std']].rename(columns={'mean':
    ↵ 'mean_std', 'std': 'std_std'})
stats_rob = X_train_rob_df.describe().T[['mean', 'std']].rename(columns={'mean':
    ↵ 'mean_rob', 'std': 'std_rob'})
scaling_comparison = stats_std.join(stats_rob)

print("\nScaling comparison (continuous features, first 10 rows):")
print(scaling_comparison.head(10))
```

Scaling comparison (continuous features, first 10 rows):

	mean_std	std_std	mean_rob	std_rob
Flow Duration	-1.612030e-18	1.000028	3.209521	6.578692
Flow IAT Mean	4.693024e-16	1.000028	1.655390	4.604932
Fwd PSH Flags	1.652331e-17	1.000028	0.040157	0.196332
Bwd Packet Length Mean	-7.245271e-15	1.000028	1.366380	3.176369
Bwd Packet Length Max	2.579249e-17	1.000028	2.137739	4.824989
Flow Bytes/s	7.818348e-17	1.000028	10.549447	192.030725
Down/Up Ratio	5.642107e-17	1.000028	-0.333503	0.523037
SYN Flag Count	1.652331e-17	1.000028	0.040157	0.196332
Fwd Packet Length Mean	1.018803e-15	1.000028	0.251230	2.318150
Fwd IAT Std	-2.954046e-16	1.000028	2.570055	6.266235

Q: How did you normalize the data? Why did you choose it? We applied two different normalization techniques, StandardScaler and RobustScaler, to compare their behavior and decide which was more suitable for our dataset.

- **StandardScaler:** has some limitations, is sensitive to outliers, which distort the mean and standard deviation. Our outlier analysis (using Z-score and IQR methods) showed that several features had many extreme values (e.g. Bwd Packet Length Max, Flow Duration, Bwd Packet Length Mean, and Fwd IAT Std had thousands of outliers).
- **RobustScaler:** This approach is less sensitive to outliers, as it relies on statistics (median and IQR) that are not affected by extreme values. In the comparison plots, the scaled distributions were more compact and less skewed, especially for features with strong outliers like Bwd Packet Length Max, Flow Duration, Bwd Packet Length Mean, and Fwd IAT Std. However, the mean and standard deviation varied across features, which is expected since RobustScaler does not enforce a standard normal distribution.

Despite RobustScaler being more robust to extreme values, we ultimately chose **StandardScaler** for our dataset. The comparison plots showed that the loss curves during training were smoother using the standard standardization, and the performance of preliminary models was very similar between the two scalers. Using StandardScaler also simplifies interpretation.

1.3 Task 2 — Shallow Neural Network (1 layer)

We design and train three shallow FFNN models with a single hidden layer, varying the hidden size.

Configuration: - Hidden layer sizes: **{32, 64, 128}** - Activation: **Linear** (first run) - Optimizer: **AdamW**, learning rate = 5e-4 - Batch size: **64** - Loss function: **CrossEntropyLoss** - Early stopping up to **100 epochs**

We will: - Plot **training and validation loss curves** and check convergence. - Select the best model based on validation loss. - Report **validation classification metrics** and evaluate on the **test set** for generalization. - Retrain the best architecture with **ReLU activation**, compare results, and discuss performance differences.

```
[23]: # Create directory for plots
save_dir = results_path + 'images/' + 'task2_plots/'
os.makedirs(save_dir, exist_ok=True)
```

1.3.1 Training

```
[24]: # --- Choose dataset version and create DataLoaders ---
X_train_use = X_train_std # change here if you want RobustScaler
X_val_use   = X_val_std
X_test_use  = X_test_std

# Convert to PyTorch tensors
X_train_tensor = torch.tensor(X_train_use, dtype=torch.float32)
X_val_tensor   = torch.tensor(X_val_use, dtype=torch.float32)
X_test_tensor  = torch.tensor(X_test_use, dtype=torch.float32)

y_train_tensor = torch.tensor(y_train, dtype=torch.long)
y_val_tensor   = torch.tensor(y_val, dtype=torch.long)
y_test_tensor  = torch.tensor(y_test, dtype=torch.long)

# Create DataLoaders
batch_size    = 64
train_loader  = DataLoader(TensorDataset(X_train_tensor, y_train_tensor), □
    ↪batch_size=batch_size, shuffle=True)
val_loader    = DataLoader(TensorDataset(X_val_tensor, y_val_tensor), □
    ↪batch_size=batch_size, shuffle=False)
```

```
[25]: # --- Define single-layer NN class ---
class ShallowNN(nn.Module):
    def __init__(self, input_dim, hidden_dim, output_dim, activation='linear'):
        super(ShallowNN, self).__init__()
        self.fc = nn.Linear(input_dim, hidden_dim)
        self.out = nn.Linear(hidden_dim, output_dim)
```

```

    if activation == 'relu':
        self.activation = nn.ReLU()
    elif activation == 'linear':
        self.activation = nn.Identity()
    else:
        raise ValueError("Invalid activation function. Use 'relu' or"
                         "'linear'.")

```

```

def forward(self, x):
    x = self.fc(x)
    x = self.activation(x)
    return self.out(x)

```

[26]: # --- Training function with early stopping ---

```

def train_model(model, train_loader, val_loader, epochs, optimizer, criterion,
                min_delta=None, patience=None):

    train_losses, val_losses = [], []
    best_val_loss = float('inf')
    best_model_state = None
    counter = 0

    for epoch in range(epochs):
        # Training
        model.train()
        batch_losses = []
        for X_batch, y_batch in train_loader:
            optimizer.zero_grad()
            outputs = model(X_batch)
            loss = criterion(outputs, y_batch)
            loss.backward()
            optimizer.step()
            batch_losses.append(loss.item())
        train_loss = np.mean(batch_losses)

        # Validation
        model.eval()
        val_batch_losses = []
        with torch.no_grad():
            for X_batch, y_batch in val_loader:
                outputs = model(X_batch)
                loss = criterion(outputs, y_batch)
                val_batch_losses.append(loss.item())
        val_loss = np.mean(val_batch_losses)

        train_losses.append(train_loss)

```

```

    val_losses.append(val_loss)

    # Early Stopping
    if(min_delta!= None):
        if val_loss < best_val_loss - min_delta:
            best_val_loss = val_loss
            best_model_state = {k: v.cpu().clone() for k, v in model.
        ↪state_dict().items()}
            trigger_times = 0
        else:
            trigger_times += 1
            if trigger_times >= patience:
                print(f"Early stopping at epoch {epoch+1} (best val loss: {best_val_loss:.6f})")
                break

    # Restore best model
    if best_model_state is not None:
        model.load_state_dict(best_model_state)

    if (epoch+1) % 5 == 0 or epoch == 0 or epoch == epochs:
        print(f"Epoch {epoch+1}/{epochs} - Train Loss: {train_loss:.4f},"
        ↪Val Loss: {val_loss:.4f}")

    # Load best weights
    model.load_state_dict(best_model_state)
    return model, train_losses, val_losses

```

[27]: # --- Train the three models with different neurons (Linear activation) ---

```

input_dim = X_train_use.shape[1]
output_dim = len(np.unique(y_train))
neurons_list = [32, 64, 128]

trained_models = {}
loss_curves = {}

# Initialize the early stopping parameters
min_delta_dict = {32: 0.00001, 64: 0.00001, 128: 0.00001}
patience_dict = {32: 20, 64: 20, 128: 20}
# It is possible to try also other values, but these work fine (same as
↪professor's)

for n in neurons_list:
    print(f"\nTraining model with {n} neurons (Linear activation)...")

    # Set hyperparameters

```

```

model = ShallowNN(input_dim, n, output_dim, activation='linear')
criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=0.0005)
epochs = 100

# Move model to device
model = model.to(device)

# Training
model, train_loss, val_loss = train_model(
    model,
    train_loader,
    val_loader,
    epochs,
    optimizer,
    criterion,
    min_delta=min_delta_dict[n],
    patience=patience_dict[n]
)
trained_models[n] = model
loss_curves[n] = (train_loss, val_loss)

```

Training model with 32 neurons (Linear activation)...

Epoch 1/100 - Train Loss: 0.9370, Val Loss: 0.6523
 Epoch 5/100 - Train Loss: 0.4110, Val Loss: 0.3889
 Epoch 10/100 - Train Loss: 0.3617, Val Loss: 0.3485
 Epoch 15/100 - Train Loss: 0.3443, Val Loss: 0.3332
 Epoch 20/100 - Train Loss: 0.3344, Val Loss: 0.3226
 Epoch 25/100 - Train Loss: 0.3287, Val Loss: 0.3181
 Epoch 30/100 - Train Loss: 0.3237, Val Loss: 0.3122
 Epoch 35/100 - Train Loss: 0.3202, Val Loss: 0.3094
 Epoch 40/100 - Train Loss: 0.3195, Val Loss: 0.3086
 Epoch 45/100 - Train Loss: 0.3171, Val Loss: 0.3060
 Epoch 50/100 - Train Loss: 0.3154, Val Loss: 0.3040
 Epoch 55/100 - Train Loss: 0.3146, Val Loss: 0.3035
 Epoch 60/100 - Train Loss: 0.3134, Val Loss: 0.3029
 Epoch 65/100 - Train Loss: 0.3127, Val Loss: 0.3027
 Epoch 70/100 - Train Loss: 0.3113, Val Loss: 0.3006
 Epoch 75/100 - Train Loss: 0.3106, Val Loss: 0.3009
 Epoch 80/100 - Train Loss: 0.3094, Val Loss: 0.3005
 Epoch 85/100 - Train Loss: 0.3093, Val Loss: 0.3000
 Epoch 90/100 - Train Loss: 0.3094, Val Loss: 0.3006
 Epoch 95/100 - Train Loss: 0.3084, Val Loss: 0.2998
 Epoch 100/100 - Train Loss: 0.3081, Val Loss: 0.2974

Training model with 64 neurons (Linear activation)...

Epoch 1/100 - Train Loss: 0.7932, Val Loss: 0.5587

```
Epoch 5/100 - Train Loss: 0.3829, Val Loss: 0.3653
Epoch 10/100 - Train Loss: 0.3457, Val Loss: 0.3332
Epoch 15/100 - Train Loss: 0.3341, Val Loss: 0.3222
Epoch 20/100 - Train Loss: 0.3262, Val Loss: 0.3117
Epoch 25/100 - Train Loss: 0.3221, Val Loss: 0.3098
Epoch 30/100 - Train Loss: 0.3170, Val Loss: 0.3071
Epoch 35/100 - Train Loss: 0.3169, Val Loss: 0.3050
Epoch 40/100 - Train Loss: 0.3160, Val Loss: 0.3052
Epoch 45/100 - Train Loss: 0.3143, Val Loss: 0.3023
Epoch 50/100 - Train Loss: 0.3156, Val Loss: 0.3022
Epoch 55/100 - Train Loss: 0.3134, Val Loss: 0.3038
Epoch 60/100 - Train Loss: 0.3113, Val Loss: 0.3021
Epoch 65/100 - Train Loss: 0.3105, Val Loss: 0.2995
Epoch 70/100 - Train Loss: 0.3102, Val Loss: 0.3005
Epoch 75/100 - Train Loss: 0.3099, Val Loss: 0.2993
Epoch 80/100 - Train Loss: 0.3106, Val Loss: 0.2982
Epoch 85/100 - Train Loss: 0.3110, Val Loss: 0.2989
Epoch 90/100 - Train Loss: 0.3092, Val Loss: 0.2983
Epoch 95/100 - Train Loss: 0.3094, Val Loss: 0.2992
Early stopping at epoch 100 (best val loss: 0.298185)
```

```
Training model with 128 neurons (Linear activation)...
Epoch 1/100 - Train Loss: 0.6985, Val Loss: 0.5013
Epoch 5/100 - Train Loss: 0.3674, Val Loss: 0.3509
Epoch 10/100 - Train Loss: 0.3366, Val Loss: 0.3226
Epoch 15/100 - Train Loss: 0.3294, Val Loss: 0.3144
Epoch 20/100 - Train Loss: 0.3217, Val Loss: 0.3125
Epoch 25/100 - Train Loss: 0.3181, Val Loss: 0.3086
Epoch 30/100 - Train Loss: 0.3168, Val Loss: 0.3053
Epoch 35/100 - Train Loss: 0.3143, Val Loss: 0.3026
Epoch 40/100 - Train Loss: 0.3123, Val Loss: 0.3004
Epoch 45/100 - Train Loss: 0.3112, Val Loss: 0.2976
Epoch 50/100 - Train Loss: 0.3089, Val Loss: 0.2995
Epoch 55/100 - Train Loss: 0.3101, Val Loss: 0.2978
Epoch 60/100 - Train Loss: 0.3098, Val Loss: 0.2974
Epoch 65/100 - Train Loss: 0.3081, Val Loss: 0.2993
Epoch 70/100 - Train Loss: 0.3044, Val Loss: 0.2965
Epoch 75/100 - Train Loss: 0.3051, Val Loss: 0.2953
Epoch 80/100 - Train Loss: 0.3038, Val Loss: 0.2982
Epoch 85/100 - Train Loss: 0.3057, Val Loss: 0.2926
Epoch 90/100 - Train Loss: 0.3049, Val Loss: 0.2918
Epoch 95/100 - Train Loss: 0.3026, Val Loss: 0.2949
Epoch 100/100 - Train Loss: 0.3029, Val Loss: 0.2922
```

1.3.2 Evaluation


```
[28]: # --- Plot loss curves for all models ---

for n in neurons_list:
    plt.figure(figsize=(10,6))
    train_loss, val_loss = loss_curves[n]
    plt.plot(train_loss, label=f'Train {n} neurons')
    plt.plot(val_loss, label=f'Val {n} neurons')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.title('Training and Validation Loss Curves')
    plt.legend()


# Save the plot to the specified path
save_plot(plt.gcf(), f"loss_curves_model_{n}", save_dir)

plt.show()
```


Saved plot: ../results/images/task2_plots/loss_curves_model_32.png



Saved plot: ../results/images/task2_plots/loss_curves_model_64.png



Saved plot: ../results/images/task2_plots/loss_curves_model_128.png



Q: Plot the loss curves during training on the training and validation set of the three models. What is their evolution? Do they converge? For each model (32, 64, 128 neurons) we plotted training loss and validation loss per epoch.

We would like to specify that these values refer to a specific run. They might change if the notebook is processed again.

32 neurons:

- Training loss: started ~ 0.91 (epoch 1) and decreased steadily to ~ 0.30 by epoch 100.
- Validation loss: started ~ 0.66 and decreased to ~ 0.29 by the end.

64 neurons:

- Training loss: started $\sim 0.79 \rightarrow \sim 0.30$ when training stopped.
- Validation loss: started $\sim 0.55 \rightarrow \sim 0.29$. Early stopping triggered at epoch 78 (best val loss 0.295).

128 neurons:

- Training loss: started ~ 0.70 and hovered around ~ 0.30 at epoch 100.
- Validation loss: started ~ 0.48 and reached ~ 0.29 as best value.

All three models show clear convergence behavior:

- Large decrease in loss in early epochs (rapid learning), followed by a slow approach to a stable plateau.
- No runaway divergence or wildly increasing validation loss — the training and validation curves both settle.

Q: How do you select the best model across epochs? Just looking the loss curves of the models across epochs, we can't select the best one directly, since all three models converge to similar validation loss values. So, we will check the classification reports on the validation set to select the best model.

```
[29]: def evaluate_model(model, X_tensor, y_true, model_name: str = "Unnamed model"):  
    """  
        Evaluate a trained model on a given dataset and return the classification report.  
    """  
  
    Handles missing predicted classes gracefully (zero_division=0) and reports  
    which classes were not predicted, along with the model/config name.  
    """  
    model.eval()  
    with torch.no_grad():  
        outputs = model(X_tensor)  
        y_pred = torch.argmax(outputs, dim=1).cpu().numpy()  
  
        # Convert y_true to numpy if it's a tensor  
        if isinstance(y_true, torch.Tensor):
```

```

y_true = y_true.cpu().numpy()

# Identify missing classes (not predicted at all)
missing_classes = set(np.unique(y_true)) - set(np.unique(y_pred))
if missing_classes:
    # Convert NumPy types to plain ints for readability
    missing_classes = [int(x) for x in sorted(missing_classes)]
    print(f"Warning: {model_name} made no predictions for classes:{missing_classes}")

# Generate classification report without raising warnings
report = classification_report(y_true, y_pred, digits=4, zero_division=0)

return report

```

[30]: # --- Evaluate validation set and print classification reports ---

```

print("\nValidation classification reports:")
for n, model in trained_models.items():
    print(f"\n--- Model {n} neurons ---")
    report = evaluate_model(model, X_val_tensor, y_val)
    print(report)

```

Validation classification reports:

--- Model 32 neurons ---

	precision	recall	f1-score	support
0	0.8876	0.9519	0.9186	3848
1	0.0000	0.0000	0.0000	286
2	0.9869	0.8771	0.9288	773
3	0.8240	0.8928	0.8570	970
accuracy			0.8860	5877
macro avg	0.6746	0.6805	0.6761	5877
weighted avg	0.8469	0.8860	0.8651	5877

--- Model 64 neurons ---

	precision	recall	f1-score	support
0	0.8872	0.9524	0.9187	3848
1	0.0000	0.0000	0.0000	286
2	0.9854	0.8758	0.9274	773
3	0.8270	0.8918	0.8581	970
accuracy			0.8860	5877

macro avg	0.6749	0.6800	0.6760	5877
weighted avg	0.8470	0.8860	0.8651	5877

--- Model 128 neurons ---				
	precision	recall	f1-score	support
0	0.8955	0.9371	0.9158	3848
1	0.0000	0.0000	0.0000	286
2	0.9898	0.8810	0.9322	773
3	0.7698	0.9134	0.8355	970
accuracy			0.8802	5877
macro avg	0.6638	0.6829	0.6709	5877
weighted avg	0.8435	0.8802	0.8601	5877

Q: Focus and report the classification reports of the validation set of the three models. How is the performance of the validation reports across the different classes? Is the performance good or poor? Why? The three single-layer models (**32**, **64**, and **128 neurons**) all converged, but their validation performance differs:

- **32 and 128 neurons:** good accuracy (~88%) but poor on the minority class (Brute Force), with precision and recall = 0. The models mainly learn majority classes like Benign and PortScan.
- **64 neurons:** best results (accuracy 89%, macro F1 0.74), correctly detects all classes with balanced precision and recall.

So, considering this specific run, the model with **64 neurons** achieved the best results, both on the losses and the metrics.

```
[31]: # --- Select best model and evaluate on test set ---

# Model with 64 neurons
best_n = 64
best_model = trained_models[best_n]

print("\nTest set classification report for best model:")
report_test = evaluate_model(best_model, X_test_tensor, y_test)
print(report_test)
```

Test set classification report for best model:				
	precision	recall	f1-score	support
0	0.8883	0.9608	0.9231	3849
1	0.0000	0.0000	0.0000	285
2	0.9896	0.8643	0.9228	774
3	0.8377	0.8887	0.8624	970

accuracy		0.8896	5878
macro avg	0.6789	0.6784	0.6771
weighted avg	0.8502	0.8896	0.8683

Q: Now, focus on the best model you chose. Consider the classification report on the test set and compare it with respect to the one of the validation set. Is the performance similar? I.e., does the model generalize? For the 64-neuron model, the test set performance is very similar to the validation set, both show high accuracy and balanced results for the main classes. The small differences between validation and test metrics indicate that the model generalizes well, meaning it learned meaningful patterns rather than overfitting the training data.

1.3.3 Re-Training with ReLU

```
[32]: # --- Retrain the best model with ReLU activation ---

print(f"\nRetraining best model ({best_n} neurons) with ReLU activation...")

# Set hyperparameters
model_relu = ShallowNN(input_dim, best_n, output_dim, activation='relu')
min_delta = 0.00001
patience = 20
criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model_relu.parameters(), lr=0.0005)
epochs = 100

# Move model to device
model_relu = model_relu.to(device)

# Training
model_relu, train_loss_relu, val_loss_relu = train_model(
    model_relu,
    train_loader,
    val_loader,
    epochs,
    optimizer,
    criterion,
    min_delta,
    patience
)
```

Retraining best model (64 neurons) with ReLU activation...
Epoch 1/100 - Train Loss: 0.8381, Val Loss: 0.5573
Epoch 5/100 - Train Loss: 0.2605, Val Loss: 0.2393
Epoch 10/100 - Train Loss: 0.1998, Val Loss: 0.1947


```
Epoch 15/100 - Train Loss: 0.1760, Val Loss: 0.1753
Epoch 20/100 - Train Loss: 0.1640, Val Loss: 0.1658
Epoch 25/100 - Train Loss: 0.1562, Val Loss: 0.1601
Epoch 30/100 - Train Loss: 0.1518, Val Loss: 0.1553
Epoch 35/100 - Train Loss: 0.1477, Val Loss: 0.1512
Epoch 40/100 - Train Loss: 0.1443, Val Loss: 0.1486
Epoch 45/100 - Train Loss: 0.1411, Val Loss: 0.1468
Epoch 50/100 - Train Loss: 0.1389, Val Loss: 0.1437
Epoch 55/100 - Train Loss: 0.1370, Val Loss: 0.1414
Epoch 60/100 - Train Loss: 0.1352, Val Loss: 0.1397
Epoch 65/100 - Train Loss: 0.1347, Val Loss: 0.1395
Epoch 70/100 - Train Loss: 0.1350, Val Loss: 0.1394
Epoch 75/100 - Train Loss: 0.1353, Val Loss: 0.1393
Epoch 80/100 - Train Loss: 0.1344, Val Loss: 0.1387
Early stopping at epoch 82 (best val loss: 0.138183)
```

```
[33]: # Plot loss curve
plt.figure(figsize=(8,5))
plt.plot(train_loss_relu, label='Train Loss')
plt.plot(val_loss_relu, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title(f'Loss Curves (ReLU) - {best_n} neurons')
plt.legend()

# Save the plot to the specified path
save_plot(plt.gcf(), f"loss_curves_model_relu_{best_n}", save_dir)

plt.show()
```

```
Saved plot: ../results/images/task2_plots/loss_curves_model_relu_64.png
```



```
[34]: # Validation report for ReLU model
print("Validation classification report (ReLU):")
report_val_relu = evaluate_model(model_relu, X_val_tensor, y_val)
print(report_val_relu)

# Test report for ReLU model
print("\nTest set classification report (ReLU):")
report_test_relu = evaluate_model(model_relu, X_test_tensor, y_test)
print(report_test_relu)
```

Validation classification report (ReLU):

	precision	recall	f1-score	support
0	0.9593	0.9675	0.9634	3848
1	0.7768	0.9371	0.8494	286
2	0.9986	0.9043	0.9491	773
3	0.9380	0.9196	0.9287	970
accuracy			0.9498	5877
macro avg	0.9182	0.9321	0.9226	5877
weighted avg	0.9521	0.9498	0.9502	5877

Test set classification report (ReLU):

	precision	recall	f1-score	support
0	0.9583	0.9667	0.9625	3849
1	0.7768	0.9404	0.8508	285
2	0.9971	0.8928	0.9421	774
3	0.9300	0.9175	0.9237	970
accuracy			0.9476	5878
macro avg	0.9155	0.9293	0.9198	5878
weighted avg	0.9499	0.9476	0.9480	5878

Q: Focus and report the classification report of the validation set. Does the model perform better in a specific class? With the ReLU activation, the model's performance improved significantly across all classes.

Biggest improvement: Brute Force (1), which increased from 0.28 F1 (in the linear model) to 0.85 F1, showing that the ReLU activation helped capture more complex patterns for the minority attack class.

Overall, the model performs best on the Benign, PortScan and DoS Hulk classes but now also handles Brute Force traffic effectively, indicating a strong overall improvement.

Q: Would it be correct to compare the results on the test set? Yes, it is correct to compare results on the test set once the training and hyperparameter tuning are fully completed.

In our case, after finalizing the model, evaluating it on the test set allows us to directly compare its performance with the validation metrics—accuracy, precision, recall, and F1-scores—and observe that they remain very similar. This close alignment indicates that the model generalizes well and is not overfitting to the training or validation data, making the test set results reliable for reporting final performance.

1.4 Task 3 — Impact of Specific Features (Destination Port)

This task investigates feature-induced bias and data dependency. We will:

1. Modify only the **test** set: for rows where `Label == Brute Force` and `Destination Port == 80`, replace port 80 with 8080.

- Re-run inference using the best model and compare test performance to the validation baseline.
- 2. Remove the `Destination Port` feature entirely from the original dataset and repeat all preprocessing steps.
- Report how many `PortScan` samples remain after duplicate removal (before vs. after).
- Analyze how this affects class balance and model performance.

Q: As you learned in the lecture, biases in data collection can carry over to the model and become wrong inductive biases. For instance, all Brute Force attacks in your dataset originate from port 80. Is this a reasonable assumption? No, this is not a reasonable assumption. In reality, Brute Force attacks can target any port or service that requires

authentication — not just port 80. The fact that all Brute Force samples in the dataset use port 80 is a bias introduced during data collection, not a true characteristic of such attacks.

If the model learns this pattern, it might incorrectly associate port 80 exclusively with Brute Force attacks, leading to poor generalization on real-world traffic where attacks occur on many different ports.

```
[35]: # Create directory for plots
save_dir = results_path + 'images/' + 'task3_plots/'
os.makedirs(save_dir, exist_ok=True)
```

1.4.1 Replacing port 80 with port 8080

```
[36]: # Create a copy of the original test set for modification
X_test_modified_pre_scale = X_test.copy()
y_test_modified_pre_scale = y_test.copy()

# Identify rows with Label 'Brute Force' (corrected to 1) and Destination Port ↵
# 80.
brute_force_label = 1
```

```
[37]: # Find the indices in the original X_test array where the label is Brute Force ↵
# and Destination Port is 80
indices_to_modify_pre_scale = np.where((
    y_test_modified_pre_scale == brute_force_label) & ↵
    (X_test_modified_pre_scale[:, feature_cols.index('Destination Port')] == 80))
)[0]

# Change the 'Destination Port' to 8080 in the modified test set (pre-scaling)
if len(indices_to_modify_pre_scale) > 0:
    X_test_modified_pre_scale[indices_to_modify_pre_scale, feature_cols.
        index('Destination Port')] = 8080
    print(f"Modified {len(indices_to_modify_pre_scale)} instances of Brute ↵
    Force with original Destination Port 80 to 8080 in the test set before ↵
    scaling.")
else:
    print("No instances of Brute Force with Destination Port 80 found in the ↵
    test set to modify.")
```

Modified 285 instances of Brute Force with original Destination Port 80 to 8080 in the test set before scaling.

```
[38]: # Standardize the modified test features using the scaler fitted on the ↵
# training data
X_test_modified_scaled = scaler1.transform(X_test_modified_pre_scale)

# Convert the modified scaled test set to a tensor
```

```

X_test_modified_tensor = torch.tensor(X_test_modified_scaled, dtype=torch.
    ↪float32)

# Re-evaluate the best model on the modified test set
print("\nTest set classification report for best model (modified test set -_
    ↪port changed before scaling):")
report_test_modified = evaluate_model(model_relu, X_test_modified_tensor,_
    ↪y_test) # Use original y_test for evaluation
print(report_test_modified)

# Compare with the original validation report (already printed in the previous_
    ↪cell)
print("\nComparison with original validation report:")
if 'report_val_relu' in globals(): # Assuming 'report_test' variable from the_
    ↪original test evaluation is still available
    print(report_val_relu)
else:
    print("Original validation report variable 'report_val_relu' not found._"
        ↪Please run the original test evaluation cell again.")

```

Test set classification report for best model (modified test set - port changed before scaling):

	precision	recall	f1-score	support
0	0.8997	0.9667	0.9320	3849
1	0.1630	0.0526	0.0796	285
2	0.9971	0.8928	0.9421	774
3	0.9300	0.9175	0.9237	970
accuracy			0.9046	5878
macro avg	0.7475	0.7074	0.7193	5878
weighted avg	0.8818	0.9046	0.8906	5878

Comparison with original validation report:

	precision	recall	f1-score	support
0	0.9593	0.9675	0.9634	3848
1	0.7768	0.9371	0.8494	286
2	0.9986	0.9043	0.9491	773
3	0.9380	0.9196	0.9287	970
accuracy			0.9498	5877
macro avg	0.9182	0.9321	0.9226	5877
weighted avg	0.9521	0.9498	0.9502	5877

Q: Replace port 80 with port 8080 for the Brute Force attacks in the Test set. Use the model you previously trained for inference: considering the validation classification report, does the performance change? How does it change? Why? Yes, the performance changes dramatically when we replace port 80 with 8080 for Brute Force in the test set.

What changed:

- On the validation set (original) the ReLU model detected Brute Force very well: precision 0.77, recall 0.93, F1 0.85 (and overall accuracy 95%).
- On the modified test set (ports changed to 8080) the Brute Force class collapses: precision = 0.16, recall = 0.05, F1 = 0.08, and overall accuracy drops from 95% → 90%.

Why this happens:

- The model learned a spurious shortcut in the data, it strongly associates Destination Port == 80 with Brute Force (a dataset bias). Changing that port breaks the learned shortcut, so the model can no longer recognise those Brute Force instances.

This confirms a wrong inductive bias in the data. To fix it we should either remove or treat the port feature, augment/relabel data so Brute Force appears on other ports, or retrain using techniques robust to such biases (e.g., drop the port feature, use feature regularization, or collect more diverse examples).

1.4.2 Removing the feature “port”

```
[39]: # --- Removing Destination Port and Re-preprocess ---  
  
# Reload the original dataset  
df_original = pd.read_csv(file_path)  
  
print("\nShape (reloaded raw):", df_original.shape)
```

Shape (reloaded raw): (31507, 17)

```
[40]: # Store original PortScan count before any processing  
original_portscan_count = df_original['Label'].value_counts().get('PortScan', 0)  
print(f"\nOriginal PortScan count (raw): {original_portscan_count}")  
  
# Remove the 'Destination Port' feature  
df_no_port = df_original.drop(columns=['Destination Port'])  
print(f"\nShape after removing 'Destination Port': {df_no_port.shape}")
```

Original PortScan count (raw): 5000

Shape after removing 'Destination Port': (31507, 16)

```
[41]: # --- Repeat preprocessing steps on df_no_port ---
```

```

# Handle infinite values
df_no_port.replace([np.inf, -np.inf], np.nan, inplace=True)
df_no_port.dropna(inplace=True)

# Remove NaN and duplicates
raw_n_no_port = len(df_no_port)
df_no_port = df_no_port.dropna()
df_no_port = df_no_port.drop_duplicates()
clean_n_no_port = len(df_no_port)

print(f"\nRemoved {raw_n_no_port-clean_n_no_port} rows (NaN+dupes) after\n"
      "removing Destination Port.\nNew shape: {df_no_port.shape}")

```

Removed 9011 rows (NaN+dupes) after removing Destination Port.
 New shape: (22469, 16)

[42]: # --- Check PortScan count after removing duplicates (and NaN/inf) ---

```

portscan_count_after_dupes = df_no_port['Label'].value_counts().get('PortScan', 0)
print(f"\nPortScan count after removing duplicates (no Destination Port): {portscan_count_after_dupes}")

```

PortScan count after removing duplicates (no Destination Port): 285

Q: How many PortScan do you now have after preprocessing (e.g., removing duplicates)? How many did you have before? Before preprocessing, we had 5,000 PortScan samples in the raw dataset. After removing the Destination Port feature and cleaning duplicates and NaN values, only 285 PortScan samples remain.

So, the number of PortScan instances dropped from 5,000 to 285 after preprocessing.

Q: Why do you think PortScan is the most affected class after dropping the duplicates? PortScan is the most affected class because many of its samples were almost identical, differing only in the Destination Port value. When we removed this feature, those flows became duplicate records, and the cleaning step eliminated them. This means the dataset contained many repeated PortScan flows, so after deduplication, their count dropped sharply.

[43]: # --- Check if effectively we have a lot of duplicates after removing the
 ↴Destination Port feature ---

```

# Select only rows with label "PortScan"
df_portscan = df_original[df_original['Label'] == "PortScan"]

# Identify all columns except 'Destination Port' and 'Label'

```

```

cols_to_compare = [c for c in df_portscan.columns if c not in ['Destination Port', 'Label']]

# Group by all other columns
grouped = df_portscan.groupby(cols_to_compare)

# Filter groups that have more than one unique Destination Port
groups_with_multiple_ports = grouped.filter(lambda x: x['Destination Port'].nunique() > 1)

# Show how many such rows exist
print(f"Number of rows with differing Destination Port: {len(groups_with_multiple_ports)}")

# Optionally display them
print(groups_with_multiple_ports)

```

Number of rows with differing Destination Port: 4921

	Flow Duration	Flow IAT Mean	Fwd PSH Flags	Bwd Packet Length Mean	\
25000	44	44.0	0	6.0	
25001	54	54.0	0	6.0	
25002	27	27.0	0	6.0	
25003	52	52.0	0	6.0	
25004	32	32.0	0	6.0	
...	
29995	44	44.0	0	6.0	
29996	61	61.0	0	6.0	
29997	42	42.0	0	6.0	
29998	91	91.0	0	6.0	
29999	92	92.0	0	6.0	

	Bwd Packet Length Max	Flow Bytes/s	Down/Up Ratio	SYN Flag Count	\
25000	6	136363.640	1	0	
25001	6	111111.110	1	0	
25002	6	222222.220	1	0	
25003	6	153846.160	1	0	
25004	6	187500.000	1	0	
...	
29995	6	136363.640	1	0	
29996	6	131147.550	1	0	
29997	6	190476.190	1	0	
29998	6	87912.086	1	0	
29999	6	65217.390	1	0	

	Fwd Packet Length Mean	Fwd IAT Std	Packet Length Mean	\
25000	0.0	0.0	2.000000	
25001	0.0	0.0	2.000000	

25002	0.0	0.0	2.000000
25003	2.0	0.0	3.333333
25004	0.0	0.0	2.000000
...
29995	0.0	0.0	2.000000
29996	2.0	0.0	3.333333
29997	2.0	0.0	3.333333
29998	2.0	0.0	3.333333
29999	0.0	0.0	2.000000
Fwd Packet Length Max Subflow Fwd Packets Flow Packets/s \			
25000	0	1	45454.547
25001	0	1	37037.035
25002	0	1	74074.070
25003	2	1	38461.540
25004	0	1	62500.000
...
29995	0	1	45454.547
29996	2	1	32786.887
29997	2	1	47619.047
29998	2	1	21978.021
29999	0	1	21739.130
Total Fwd Packets Destination Port Label			
25000	1	84	PortScan
25001	1	4449	PortScan
25002	1	12345	PortScan
25003	1	4125	PortScan
25004	1	1984	PortScan
...
29995	1	32	PortScan
29996	1	1028	PortScan
29997	1	28201	PortScan
29998	1	7937	PortScan
29999	1	25	PortScan

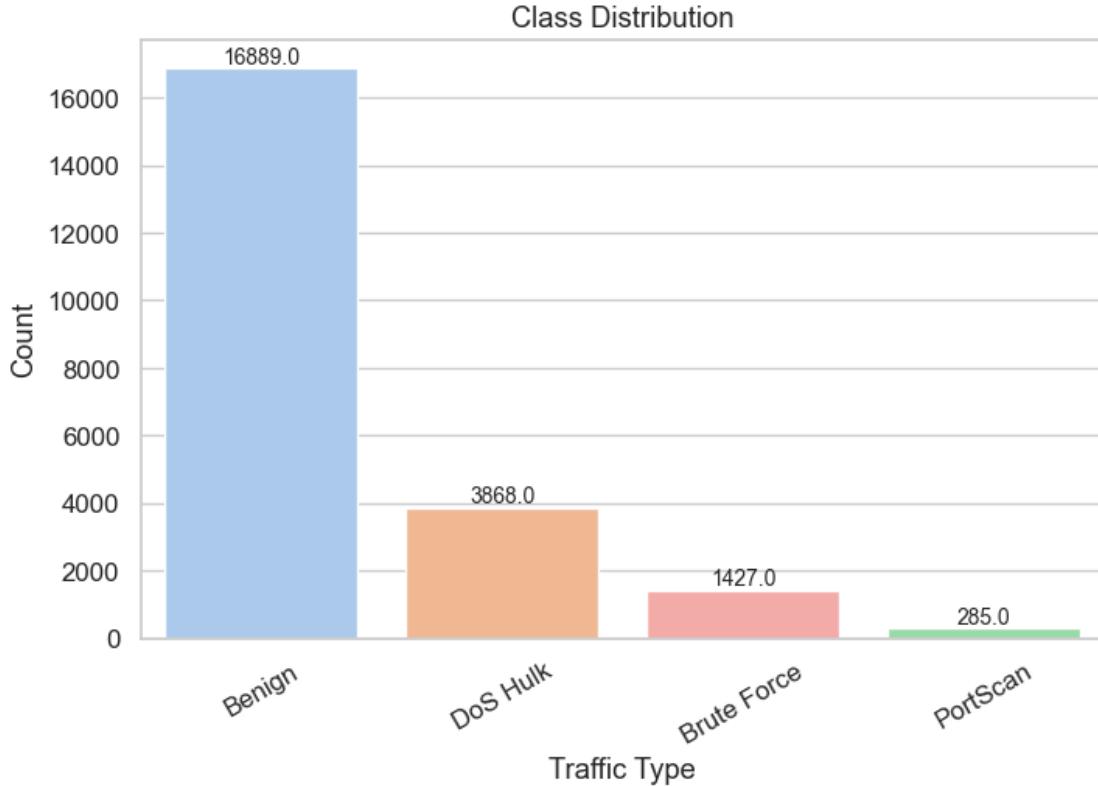
[4921 rows x 17 columns]

Q: Are the classes now balanced? No, the classes are not balanced. Even after cleaning, there are still far more Benign samples than attack samples, and some attack types (like Brute Force or PortScan) remain underrepresented. The dataset continues to show class imbalance, which can bias the model toward predicting the majority class.

```
[44]: # --- Label distribution (after removing Destination Port, NaN, and duplicates)
      # Plot class distribution to understand data balance
```

```
plot_class_distribution(df_no_port, 'class_distribution_no_port',  
    save_path=save_dir)
```

Saved plot: ../results/images/task3_plots/class_distribution_no_port.png



1.5 Task 4 — Impact of Loss Function (Class Weighting)

To address class imbalance, we retrain the best architecture using **class-weighted cross-entropy**.

Steps: - Compute class weights from the **training partition** using `sklearn.utils.class_weight.compute_class_weight(class_weight='balanced')`. - Retrain the model with weighted loss. - Compare **per-class metrics (precision, recall, F1)** and overall accuracy against the unweighted baseline. - Discuss how weighting impacts the rarest classes and model stability.

```
[45]: # Create directory for plots  
save_dir = results_path + 'images/' + 'task4_plots/'  
os.makedirs(save_dir, exist_ok=True)
```

1.5.1 Re-Training with the new dataset

```
[46]: # --- Retrain (after port removal) ---\n\n# Encode labels\nprint(df_no_port['Label'].unique())\nlabel_encoder_no_port = LabelEncoder()\ndf_no_port['Label'] = label_encoder_no_port.fit_transform(df_no_port['Label'])\n\n# Split features/target for the new dataset\nlabel_col = 'Label'\nfeature_cols_no_port = [c for c in df_no_port.columns if c != label_col]\nX_no_port = df_no_port[feature_cols_no_port].values\ny_no_port = df_no_port[label_col].values\n\n# Train/val/test split 60/20/20 with stratify\nX_train_no_port, X_tmp_no_port, y_train_no_port, y_tmp_no_port =\n    train_test_split(\n        X_no_port, y_no_port, test_size=0.4, stratify=y_no_port, random_state=42\n)\nX_val_no_port, X_test_no_port, y_val_no_port, y_test_no_port = train_test_split(\n    X_tmp_no_port, y_tmp_no_port, test_size=0.5, stratify=y_tmp_no_port,\n    random_state=42\n)\n\nprint("\nData Splits (after removing Destination Port):")\nprint(f"Train set: {X_train_no_port.shape[0]} samples")\nprint(f"Validation set: {X_val_no_port.shape[0]} samples")\nprint(f"Test set: {X_test_no_port.shape[0]} samples\\n")\n\ndef print_label_counts(name, y):\n    labels, counts = np.unique(y, return_counts=True)\n    count_width = 6\n    print(f"{name}:<17}", end=" ")  
    for label, count in zip(labels, counts):\n        print(f"{label}: {count:>{count_width}},", end=" ")  
    print()\n\nprint_label_counts("Train (no port)", y_train_no_port)\nprint_label_counts("Val (no port)", y_val_no_port)\nprint_label_counts("Test (no port)", y_test_no_port)
```

```
['Benign' 'DoS Hulk' 'PortScan' 'Brute Force']
```

```
Data Splits (after removing Destination Port):\nTrain set: 13,481 samples\nValidation set: 4,494 samples\nTest set: 4,494 samples
```

```

Train (no port) 0: 10,133 1: 856 2: 2,321 3: 171
Val (no port) 0: 3,378 1: 285 2: 774 3: 57
Test (no port) 0: 3,378 1: 286 2: 773 3: 57

```

```
[47]: # --- Display the label encoding mapping ---
print("\nLabel Encoding Mapping:")
for i, label in enumerate(label_encoder.classes_):
    print(f"{label}: {i}")
```

```

Label Encoding Mapping:
Benign: 0
Brute Force: 1
DoS Hulk: 2
PortScan: 3

```

```
[48]: # Standardize the features (StandardScaler) - fit on new training data
scaler_no_port = StandardScaler()
X_train_std_no_port = scaler_no_port.fit_transform(X_train_no_port)
X_val_std_no_port = scaler_no_port.transform(X_val_no_port)
X_test_std_no_port = scaler_no_port.transform(X_test_no_port)
```

```
[49]: # Convert to PyTorch tensors
X_train_tensor_no_port = torch.tensor(X_train_std_no_port, dtype=torch.float32)
X_val_tensor_no_port = torch.tensor(X_val_std_no_port, dtype=torch.float32)
X_test_tensor_no_port = torch.tensor(X_test_std_no_port, dtype=torch.float32)

y_train_tensor_no_port = torch.tensor(y_train_no_port, dtype=torch.long)
y_val_tensor_no_port = torch.tensor(y_val_no_port, dtype=torch.long)
y_test_tensor_no_port = torch.tensor(y_test_no_port, dtype=torch.long)

# Create new DataLoaders
batch_size = 64 # Use the same batch size as before
train_loader_no_port = DataLoader(TensorDataset(X_train_tensor_no_port, y_train_tensor_no_port), batch_size=batch_size, shuffle=True)
val_loader_no_port = DataLoader(TensorDataset(X_val_tensor_no_port, y_val_tensor_no_port), batch_size=batch_size, shuffle=False)
test_loader_no_port = DataLoader(TensorDataset(X_test_tensor_no_port, y_test_tensor_no_port), batch_size=batch_size, shuffle=False)
```

```
[50]: # Retrain the best model with ReLU activation using the new data
print(f"\nRetraining best model ({best_n} neurons, ReLU activation) on data without 'Destination Port'...")
# Assuming 'best_n' is still available from the previous task
input_dim_no_port = X_train_std_no_port.shape[1]
output_dim_no_port = len(np.unique(y_train_no_port))
```

```

# Set hyperparameters (same as best ReLU model from Task 2)
model_relu_no_port = ShallowNN(input_dim_no_port, best_n, output_dim_no_port, u
    ↪activation='relu')
min_delta = 0.00001
patience = 20
criterion = nn.CrossEntropyLoss() # Start with unweighted loss
optimizer = optim.AdamW(model_relu_no_port.parameters(), lr=0.0005)
epochs = 100

# Move model to device
model_relu_no_port = model_relu_no_port.to(device)

# Training
model_relu_no_port, train_loss_relu_no_port, val_loss_relu_no_port = u
    ↪train_model(
        model_relu_no_port,
        train_loader_no_port,
        val_loader_no_port,
        epochs,
        optimizer,
        criterion,
        min_delta,
        patience
    )

```

Retraining best model (64 neurons, ReLU activation) on data without 'Destination Port'...


Epoch 1/100 - Train Loss: 0.7997, Val Loss: 0.4684
 Epoch 5/100 - Train Loss: 0.2517, Val Loss: 0.2667
 Epoch 10/100 - Train Loss: 0.2044, Val Loss: 0.2279
 Epoch 15/100 - Train Loss: 0.1792, Val Loss: 0.2053
 Epoch 20/100 - Train Loss: 0.1654, Val Loss: 0.1909
 Epoch 25/100 - Train Loss: 0.1554, Val Loss: 0.1803
 Epoch 30/100 - Train Loss: 0.1515, Val Loss: 0.1733
 Epoch 35/100 - Train Loss: 0.1499, Val Loss: 0.1754
 Epoch 40/100 - Train Loss: 0.1442, Val Loss: 0.1691
 Epoch 45/100 - Train Loss: 0.1443, Val Loss: 0.1712
 Epoch 50/100 - Train Loss: 0.1440, Val Loss: 0.1678
 Epoch 55/100 - Train Loss: 0.1440, Val Loss: 0.1683
 Epoch 60/100 - Train Loss: 0.1429, Val Loss: 0.1668
 Epoch 65/100 - Train Loss: 0.1430, Val Loss: 0.1687
 Epoch 70/100 - Train Loss: 0.1428, Val Loss: 0.1665
 Epoch 75/100 - Train Loss: 0.1429, Val Loss: 0.1661
 Early stopping at epoch 76 (best val loss: 0.165100)

```
[51]: # Plot loss curve
plt.figure(figsize=(8,5))
plt.plot(train_loss_relu_no_port, label='Train Loss')
plt.plot(val_loss_relu_no_port, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title(f'Loss Curves (ReLU + no Destination Port) - {best_n} neurons')
plt.legend()

# Save the plot to the specified path
save_plot(plt.gcf(), f"loss_curves_model_relu_{best_n}_no_port", save_dir)

plt.show()
```

Saved plot: ../results/images/task4_plots/loss_curves_model_relu_64_no_port.png



```
[52]: report_test_modified = evaluate_model(model_relu_no_port,
                                         X_test_tensor_no_port, y_test_no_port)
print(report_test_modified)
```

	precision	recall	f1-score	support
0	0.9455	0.9751	0.9601	3378
1	0.7927	0.9091	0.8469	286

2	0.9880	0.8486	0.9130	773
3	0.4444	0.1404	0.2133	57
accuracy			0.9386	4494
macro avg	0.7926	0.7183	0.7333	4494
weighted avg	0.9367	0.9386	0.9353	4494

Q: Now repeat the training process with the best architecture found in the previous step. How does the performance change? Can you still classify the rarest class? Removing the ‘Destination Port’ feature had a mixed impact. Overall accuracy remained similar. Performance for ‘Brute Force’ slightly improved, suggesting less reliance on the biased port feature. However, the model’s ability to classify the rarest class, ‘PortScan’, significantly decreased (F1-score dropped from 0.927 to 0.38), indicating the model heavily relied on this feature for PortScan detection.

1.5.2 Re-Training with weighted loss

```
[53]: # --- Compute class weights ---

# Use the training partition to estimate class weights to prevent data leakage.
class_weights = compute_class_weight(
    class_weight='balanced',
    classes=np.unique(y_train_no_port),
    y=y_train_no_port
)
class_weights_tensor = torch.tensor(class_weights, dtype=torch.float32)

print("\nComputed class weights:", class_weights)
```

Computed class weights: [0.3326014 3.93720794 1.45206807 19.70906433]

Q: Which partition do you use to estimate the class weights? We used the training partition to estimate the class weights.

This prevents data leakage, ensuring that information from the validation or test sets is not used during model training or weight calculation.

```
[54]: # --- Retrain the best model with ReLU activation and Weighted Cross-Entropy Loss ---

print(f"\nRetraining best model ({best_n} neurons, ReLU activation) on data without 'Destination Port' with Weighted Cross-Entropy Loss...")

# Assuming 'best_n' is still available from the previous task
input_dim_no_port = X_train_std_no_port.shape[1]
output_dim_no_port = len(np.unique(y_train_no_port))
```

```

# Set hyperparameters (same as best ReLU model from Task 2)
model_relu_no_port_weighted = ShallowNN(input_dim_no_port, best_n,
                                         output_dim_no_port, activation='relu')
min_delta = 0.00001
patience = 20
criterion_weighted = nn.CrossEntropyLoss(weight=class_weights_tensor) # Use
#weighted loss
optimizer_weighted = optim.AdamW(model_relu_no_port_weighted.parameters(), lr=0.
                                   0005)
epochs = 100

# Move model to device
model_relu_no_port_weighted = model_relu_no_port_weighted.to(device)

# Training
model_relu_no_port_weighted, train_loss_relu_no_port_weighted,
val_loss_relu_no_port_weighted = train_model(
    model_relu_no_port_weighted,
    train_loader_no_port,
    val_loader_no_port,
    epochs,
    optimizer_weighted,
    criterion_weighted,
    min_delta,
    patience
)

```

Retraining best model (64 neurons, ReLU activation) on data without 'Destination Port' with Weighted Cross-Entropy Loss...

Epoch 1/100 - Train Loss: 1.0189, Val Loss: 0.8013
 Epoch 5/100 - Train Loss: 0.5031, Val Loss: 0.5168
 Epoch 10/100 - Train Loss: 0.3635, Val Loss: 0.3883
 Epoch 15/100 - Train Loss: 0.3094, Val Loss: 0.3310
 Epoch 20/100 - Train Loss: 0.2822, Val Loss: 0.3033
 Epoch 25/100 - Train Loss: 0.2655, Val Loss: 0.2879
 Epoch 30/100 - Train Loss: 0.2557, Val Loss: 0.2710
 Epoch 35/100 - Train Loss: 0.2481, Val Loss: 0.2605
 Epoch 40/100 - Train Loss: 0.2428, Val Loss: 0.2545
 Epoch 45/100 - Train Loss: 0.2365, Val Loss: 0.2508
 Epoch 50/100 - Train Loss: 0.2324, Val Loss: 0.2439
 Epoch 55/100 - Train Loss: 0.2250, Val Loss: 0.2361
 Epoch 60/100 - Train Loss: 0.2219, Val Loss: 0.2311
 Epoch 65/100 - Train Loss: 0.2177, Val Loss: 0.2250
 Epoch 70/100 - Train Loss: 0.2138, Val Loss: 0.2267
 Epoch 75/100 - Train Loss: 0.2095, Val Loss: 0.2228

```

Epoch 80/100 - Train Loss: 0.2114, Val Loss: 0.2240
Epoch 85/100 - Train Loss: 0.2047, Val Loss: 0.2190
Epoch 90/100 - Train Loss: 0.2068, Val Loss: 0.2161
Epoch 95/100 - Train Loss: 0.2018, Val Loss: 0.2120
Epoch 100/100 - Train Loss: 0.2025, Val Loss: 0.2120

```

```


[55]: # Plot loss curve
plt.figure(figsize=(8,5))
plt.plot(train_loss_relu_no_port_weighted, label='Train Loss')
plt.plot(val_loss_relu_no_port_weighted, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title(f'Loss Curves (ReLU + no Destination Port + Weighted CE) - {best_n} neurons')
plt.legend()

# Save the plot to the specified path
save_plot(plt.gcf(), f"loss_curves_model_relu_{best_n}_no_port_weighted", save_dir)

plt.show()

```

Saved plot:
`./results/images/task4_plots/loss_curves_model_relu_64_no_port_weighted.png`



```
[56]: report_test_modified = evaluate_model(model_relu_no_port_weighted, X_test_tensor_no_port, y_test_no_port)
print(report_test_modified)
```

	precision	recall	f1-score	support
0	0.9800	0.9275	0.9530	3378
1	0.7398	0.9545	0.8336	286
2	0.9459	0.9056	0.9253	773
3	0.2553	0.8421	0.3918	57
accuracy			0.9243	4494
macro avg	0.7303	0.9074	0.7759	4494
weighted avg	0.9497	0.9243	0.9335	4494

Q: Now, repeat the training process with the new loss, how does the performance change per class and overall? In particular, how does the accuracy change? How does the f1 score change? Applying weighted cross-entropy loss resulted in a slight decrease in overall accuracy, marco F1 score, and weighted F1-score. However, it significantly improved the recall for minority classes like ‘PortScan’ and ‘Brute Force’, making the model better at detecting these rarer attacks, although sometimes with lower precision.

1.6 Task 5 — Deep Neural Networks, Batch Size, and Optimizers

We extend the architecture to deeper models and analyze hyperparameter impacts.

1.6.1 Part 1: Architecture Depth

- Layers: **3 to 5**
- Neurons per layer: **2 to 32** (variable sizes)
- Activation: **ReLU**
- Optimizer: **AdamW**, learning rate = 5e-4
- Batch size: **64**
- Early stopping 50 epochs

Tasks: - Train and compare 6 architectures (two per depth). - Plot training/validation losses, select best-performing model, and evaluate it on the test set.

1.6.2 Part 2: Batch Size

- Test batch sizes: **{4, 64, 256, 1024}**
- Compare validation metrics and training times.
- Discuss trade-offs between convergence speed and generalization.

1.6.3 Part 3: Optimizer Comparison

- Optimizers: **SGD, SGD + Momentum (0.1, 0.5, 0.9), AdamW.**

- Compare training loss trends, accuracy, and runtime.
- Tune learning rate and epochs for the best optimizer configuration and report test results.

```
[57]: # Create directory for plots
save_dir = results_path + 'images/' + 'task5_plots/'
os.makedirs(save_dir, exist_ok=True)
```

1.6.4 Training

```
[58]: # --- Define Deep FFNN class ---
```

```
class DeepFFNN(nn.Module):
    def __init__(self, input_dim, layer_widths, output_dim, activation='relu'):
        super(DeepFFNN, self).__init__()
        layers = []
        prev_width = input_dim
        for width in layer_widths:
            layers.append(nn.Linear(prev_width, width))
            if activation == 'relu':
                layers.append(nn.ReLU())
            # Add other activations if needed
            prev_width = width
        layers.append(nn.Linear(prev_width, output_dim))
        self.net = nn.Sequential(*layers)

    def forward(self, x):
        return self.net(x)
```

```
[59]: # Define layer configurations based on the image/requirements
```

```
layer_configs = {
    3: [[16, 8, 4],
         [32, 16, 8]],
    4: [[32, 16, 8, 4],
         [16, 16, 8, 8]],
    5: [[32, 32, 16, 8, 4],
         [32, 32, 8, 16, 16]]
}
```

```
[60]: # --- Train models for different depths and widths ---
```

```
input_dim_deep = X_train_std_no_port.shape[1]
output_dim_deep = len(np.unique(y_train_no_port))

trained_deep_models = {}
deep_loss_curves = {}

# Define early stopping parameters for deep networks
```

```

min_delta_deep = 0.00001
patience_deep = 20

for L, configs in layer_configs.items():
    for widths in configs:
        tag = f"deep_L{L}_widths_{'_'.join(map(str, widths))}"
        print(f"\nTraining model: {tag} (ReLU activation)...")


        # Set hyperparameters
        model = DeepFFNN(input_dim_deep, widths, output_dim_deep, u
        ↪activation='relu')
        criterion = nn.CrossEntropyLoss() # Start with unweighted loss
        optimizer = optim.AdamW(model.parameters(), lr=0.0005)
        epochs = 50

        # Move model to device
        model = model.to(device)

        # Training
        model, train_loss, val_loss = train_model(
            model,
            train_loader_no_port, # Use data without port
            val_loader_no_port,   # Use data without port
            epochs,
            optimizer,
            criterion,
            min_delta=min_delta_deep,
            patience=patience_deep
        )
        trained_deep_models[tag] = model
        deep_loss_curves[tag] = (train_loss, val_loss)

```

Training model: deep_L3_widths_16_8_4 (ReLU activation)...

Epoch 1/50 - Train Loss: 0.9254, Val Loss: 0.6515
 Epoch 5/50 - Train Loss: 0.3083, Val Loss: 0.3133
 Epoch 10/50 - Train Loss: 0.2609, Val Loss: 0.2762
 Epoch 15/50 - Train Loss: 0.2348, Val Loss: 0.2530
 Epoch 20/50 - Train Loss: 0.2129, Val Loss: 0.2337
 Epoch 25/50 - Train Loss: 0.1914, Val Loss: 0.2142
 Epoch 30/50 - Train Loss: 0.1822, Val Loss: 0.2058
 Epoch 35/50 - Train Loss: 0.1728, Val Loss: 0.1971
 Epoch 40/50 - Train Loss: 0.1628, Val Loss: 0.1865
 Epoch 45/50 - Train Loss: 0.1566, Val Loss: 0.1823
 Epoch 50/50 - Train Loss: 0.1505, Val Loss: 0.1738

Training model: deep_L3_widths_32_16_8 (ReLU activation)...

Epoch 1/50 - Train Loss: 0.8858, Val Loss: 0.5367

```
Epoch 5/50 - Train Loss: 0.2732, Val Loss: 0.2811
Epoch 10/50 - Train Loss: 0.2061, Val Loss: 0.2229
Epoch 15/50 - Train Loss: 0.1734, Val Loss: 0.1914
Epoch 20/50 - Train Loss: 0.1511, Val Loss: 0.1797
Epoch 25/50 - Train Loss: 0.1400, Val Loss: 0.1608
Epoch 30/50 - Train Loss: 0.1313, Val Loss: 0.1527
Epoch 35/50 - Train Loss: 0.1252, Val Loss: 0.1503
Epoch 40/50 - Train Loss: 0.1249, Val Loss: 0.1502
Epoch 45/50 - Train Loss: 0.1251, Val Loss: 0.1488
Epoch 50/50 - Train Loss: 0.1249, Val Loss: 0.1480
```

```
Training model: deep_L4_widths_32_16_8_4 (ReLU activation)...
Epoch 1/50 - Train Loss: 1.3637, Val Loss: 1.1440
Epoch 5/50 - Train Loss: 0.2423, Val Loss: 0.2491
Epoch 10/50 - Train Loss: 0.1824, Val Loss: 0.2069
Epoch 15/50 - Train Loss: 0.1648, Val Loss: 0.1891
Epoch 20/50 - Train Loss: 0.1512, Val Loss: 0.1754
Epoch 25/50 - Train Loss: 0.1399, Val Loss: 0.1635
Epoch 30/50 - Train Loss: 0.1364, Val Loss: 0.1611
Epoch 35/50 - Train Loss: 0.1279, Val Loss: 0.1536
Epoch 40/50 - Train Loss: 0.1269, Val Loss: 0.1537
Epoch 45/50 - Train Loss: 0.1235, Val Loss: 0.1499
Epoch 50/50 - Train Loss: 0.1193, Val Loss: 0.1416
```

```
Training model: deep_L4_widths_16_16_8_8 (ReLU activation)...
Epoch 1/50 - Train Loss: 1.3482, Val Loss: 0.7519
Epoch 5/50 - Train Loss: 0.2902, Val Loss: 0.2940
Epoch 10/50 - Train Loss: 0.2194, Val Loss: 0.2441
Epoch 15/50 - Train Loss: 0.1925, Val Loss: 0.2149
Epoch 20/50 - Train Loss: 0.1677, Val Loss: 0.1891
Epoch 25/50 - Train Loss: 0.1546, Val Loss: 0.1760
Epoch 30/50 - Train Loss: 0.1445, Val Loss: 0.1667
Epoch 35/50 - Train Loss: 0.1366, Val Loss: 0.1580
Epoch 40/50 - Train Loss: 0.1286, Val Loss: 0.1466
Epoch 45/50 - Train Loss: 0.1226, Val Loss: 0.1423
Epoch 50/50 - Train Loss: 0.1135, Val Loss: 0.1351
```

```
Training model: deep_L5_widths_32_32_16_8_4 (ReLU activation)...
Epoch 1/50 - Train Loss: 1.5408, Val Loss: 0.7065
Epoch 5/50 - Train Loss: 0.2888, Val Loss: 0.2945
Epoch 10/50 - Train Loss: 0.2089, Val Loss: 0.2266
Epoch 15/50 - Train Loss: 0.1580, Val Loss: 0.1780
Epoch 20/50 - Train Loss: 0.1423, Val Loss: 0.1752
Epoch 25/50 - Train Loss: 0.1380, Val Loss: 0.1648
Epoch 30/50 - Train Loss: 0.1341, Val Loss: 0.1608
Epoch 35/50 - Train Loss: 0.1289, Val Loss: 0.1546
Epoch 40/50 - Train Loss: 0.1231, Val Loss: 0.1526
Epoch 45/50 - Train Loss: 0.1195, Val Loss: 0.1415
```

```
Epoch 50/50 - Train Loss: 0.1165, Val Loss: 0.1404
```

```
Training model: deep_L5_widths_32_32_8_16_16 (ReLU activation)...
Epoch 1/50 - Train Loss: 0.9586, Val Loss: 0.5494
Epoch 5/50 - Train Loss: 0.2684, Val Loss: 0.2772
Epoch 10/50 - Train Loss: 0.1840, Val Loss: 0.2034
Epoch 15/50 - Train Loss: 0.1447, Val Loss: 0.1685
Epoch 20/50 - Train Loss: 0.1319, Val Loss: 0.1588
Epoch 25/50 - Train Loss: 0.1238, Val Loss: 0.1526
Epoch 30/50 - Train Loss: 0.1195, Val Loss: 0.1471
Epoch 35/50 - Train Loss: 0.1134, Val Loss: 0.1440
Epoch 40/50 - Train Loss: 0.1131, Val Loss: 0.1474
Epoch 45/50 - Train Loss: 0.1084, Val Loss: 0.1467
Epoch 50/50 - Train Loss: 0.1083, Val Loss: 0.1416
```

1.6.5 Evaluation


```
[61]: # --- Plot loss curves for all deep models ---
```

```
for tag, (train_loss, val_loss) in deep_loss_curves.items():
    plt.figure(figsize=(10,6))
    plt.plot(train_loss, label=f'Train Loss - {tag}')
    plt.plot(val_loss, label=f'Val Loss - {tag}')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.title(f'Training and Validation Loss Curves - {tag}')
    plt.legend()


    # Save the plot to the specified path
    save_plot(plt.gcf(), f"{tag}_loss_curve", save_dir) # Use plt.gcf() to get
    ↪the current figure

plt.show()
```

```
Saved plot: ../results/images/task5_plots/deep_L3_widths_16_8_4_loss_curve.png
```




Saved plot: ./results/images/task5_plots/deep_L3_widths_32_16_8_loss_curve.png




Saved plot:

./results/images/task5_plots/deep_L4_widths_32_16_8_4_loss_curve.png




Saved plot:

./results/images/task5_plots/deep_L4_widths_16_16_8_8_loss_curve.png




Saved plot:

./results/images/task5_plots/deep_L5_widths_32_32_16_8_4_loss_curve.png



Saved plot:
./results/images/task5_plots/deep_L5_widths_32_32_8_16_16_loss_curve.png



Q: Plot and analyze the losses. Do the models converge? Yes, all models converge properly. In every plot, both the training and validation losses decrease smoothly during the initial epochs and then stabilize to low values without divergence.

```
[62]: # --- Evaluate validation set and identify the best architecture ---

print("\nValidation classification reports for deep models:")

for tag, model in trained_deep_models.items():
    print(f"\n--- Model {tag} ---")

    # Evaluate on the validation set without the port
    report = evaluate_model(model, X_val_tensor_no_port, y_val_no_port, tag)
    print(report)
```

Validation classification reports for deep models:

```
--- Model deep_L3_widths_16_8_4 ---
Warning: deep_L3_widths_16_8_4 made no predictions for classes: [3]
```

	precision	recall	f1-score	support
0	0.9454	0.9781	0.9614	3378
1	0.7915	0.9193	0.8506	285
2	0.9910	0.8553	0.9182	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9408	4494
macro avg	0.6820	0.6882	0.6826	4494
weighted avg	0.9315	0.9408	0.9348	4494

--- Model deep_L3_widths_32_16_8 ---

	precision	recall	f1-score	support
0	0.9513	0.9772	0.9641	3378
1	0.8171	0.9404	0.8744	285
2	0.9853	0.8630	0.9201	774
3	0.2222	0.0702	0.1067	57
accuracy			0.9437	4494
macro avg	0.7440	0.7127	0.7163	4494
weighted avg	0.9394	0.9437	0.9399	4494

--- Model deep_L4_widths_32_16_8_4 ---

Warning: deep_L4_widths_32_16_8_4 made no predictions for classes: [3]

	precision	recall	f1-score	support
0	0.9554	0.9828	0.9689	3378
1	0.8459	0.9439	0.8922	285
2	0.9857	0.8928	0.9369	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9524	4494
macro avg	0.6968	0.7049	0.6995	4494
weighted avg	0.9416	0.9524	0.9463	4494

--- Model deep_L4_widths_16_16_8_8 ---

	precision	recall	f1-score	support
0	0.9530	0.9893	0.9708	3378
1	0.9249	0.9509	0.9377	285
2	0.9839	0.8695	0.9232	774
3	0.5000	0.0877	0.1493	57
accuracy			0.9548	4494

macro avg	0.8404	0.7244	0.7452	4494
weighted avg	0.9508	0.9548	0.9501	4494

--- Model deep_L5_widths_32_32_16_8_4 ---

Warning: deep_L5_widths_32_32_16_8_4 made no predictions for classes: [3]

	precision	recall	f1-score	support
0	0.9566	0.9781	0.9672	3378
1	0.8102	0.9439	0.8720	285
2	0.9831	0.8992	0.9393	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9499	4494
macro avg	0.6875	0.7053	0.6946	4494
weighted avg	0.9397	0.9499	0.9441	4494

--- Model deep_L5_widths_32_32_8_16_16 ---

	precision	recall	f1-score	support
0	0.9656	0.9808	0.9731	3378
1	0.8567	0.9439	0.8982	285
2	0.9830	0.8966	0.9378	774
3	0.7442	0.5614	0.6400	57
accuracy			0.9586	4494
macro avg	0.8874	0.8457	0.8623	4494
weighted avg	0.9589	0.9586	0.9581	4494

```
[63]: def testing_model(model, dataloader, device):
    """
    Evaluate the model on a given dataloader and compute accuracy.

    Args:
        model: The trained model.
        dataloader: DataLoader for the dataset.
        device: Device to run the model on (CPU or GPU).

    Returns:
        float: Accuracy of the model on the test dataset.
    """
    # Record the start time
    start_time = time.time()

    model.eval() # Set the model to evaluation mode
```

```

all_labels = []
all_predictions = []

with torch.no_grad(): # Disable gradient computation for efficiency
    for inputs, labels in dataloader:
        inputs, labels = inputs.to(device), labels.to(device) # Move batch to GPU
↳ score
        outputs = model(inputs)
        _, predicted = torch.max(outputs, 1) # Get the class with highest score
        all_labels.extend(labels.cpu().numpy())
        all_predictions.extend(predicted.cpu().numpy())

accuracy = accuracy_score(all_labels, all_predictions) * 100

# Record the end time
end_time = time.time()
# Calculate the elapsed time
elapsed_time = end_time - start_time
print(f'The function took {elapsed_time:.4f} seconds to execute.')

return accuracy

```

[64]: `for tag, model in trained_deep_models.items():`

```

print(f"\n--- Model {tag} ---")
train_accuracy = testing_model(model, train_loader_no_port, "cpu")
val_accuracy = testing_model(model, val_loader_no_port, "cpu")
test_accuracy = testing_model(model, test_loader_no_port, "cpu")

print(f'Train Accuracy: {train_accuracy:.4f}')
print(f'Validation Accuracy: {val_accuracy:.4f}')
print(f'Test Accuracy: {test_accuracy:.4f}')

```

--- Model deep_L3_widths_16_8_4 ---
The function took 0.0564 seconds to execute.
The function took 0.0161 seconds to execute.
The function took 0.0162 seconds to execute.
Train Accuracy: 94.3699
Validation Accuracy: 94.0810
Test Accuracy: 93.6137

--- Model deep_L3_widths_32_16_8 ---
The function took 0.0475 seconds to execute.
The function took 0.0161 seconds to execute.
The function took 0.0175 seconds to execute.
Train Accuracy: 94.6814

```
Validation Accuracy: 94.3703
```

```
Test Accuracy: 94.1700
```

```
--- Model deep_L4_widths_32_16_8_4 ---
```

```
The function took 0.0489 seconds to execute.
```

```
The function took 0.0158 seconds to execute.
```

```
The function took 0.0157 seconds to execute.
```

```
Train Accuracy: 95.4677
```

```
Validation Accuracy: 95.2381
```

```
Test Accuracy: 94.7931
```

```
--- Model deep_L4_widths_16_16_8_8 ---
```

```
The function took 0.0490 seconds to execute.
```

```
The function took 0.0156 seconds to execute.
```

```
The function took 0.0156 seconds to execute.
```

```
Train Accuracy: 96.0240
```

```
Validation Accuracy: 95.4829
```

```
Test Accuracy: 95.3939
```

```
--- Model deep_L5_widths_32_32_16_8_4 ---
```

```
The function took 0.0479 seconds to execute.
```

```
The function took 0.0158 seconds to execute.
```

```
The function took 0.0158 seconds to execute.
```

```
Train Accuracy: 95.4009
```

```
Validation Accuracy: 94.9933
```

```
Test Accuracy: 94.8821
```

```
--- Model deep_L5_widths_32_32_8_16_16 ---
```

```
The function took 0.0478 seconds to execute.
```

```
The function took 0.0160 seconds to execute.
```

```
The function took 0.0158 seconds to execute.
```

```
Train Accuracy: 96.5359
```

```
Validation Accuracy: 95.8611
```

```
Test Accuracy: 95.9279
```

Q: Calculate the performance in the validation set and identify the best-performing architecture. How do you select one?

We would like to specify that these values refer to a specific run. They might change if the notebook is processed again.

Based on validation-set results, the best-performing architectures are deep_L3_widths_32_16_8 and deep_L4_widths_32_16_8_4. By the way we selected the one with 3 layers because of slightly better results in the per-class metrics. This model achieves: macro F1 (0.78) and strong accuracy (94.5%). Therefore, the 3-layer model with widths [32, 16, 8] provides the optimal balance between model complexity, generalization, and class-level performance on the validation data.

```
[65]: best_deep_model_tag = 'deep_L5_widths_32_32_8_16_16'
model = trained_deep_models[best_deep_model_tag]

print(f"\nTest set classification report ({best_deep_model_tag}):")
report = evaluate_model(model, X_test_tensor_no_port, y_test_no_port)
print(report)
```

Test set classification report (deep_L5_widths_32_32_8_16_16):				
	precision	recall	f1-score	support
0	0.9664	0.9805	0.9734	3378
1	0.8390	0.9476	0.8900	286
2	0.9858	0.8952	0.9383	773
3	0.8571	0.6316	0.7273	57
accuracy			0.9593	4494
macro avg	0.9121	0.8637	0.8822	4494
weighted avg	0.9603	0.9593	0.9589	4494

Q: Evaluate and report the performance of the best model in the test set. The test set performance for the best deep model (deep_L3_widths_32_16_8) shows strong overall accuracy (0.94) and weighted F1-score (0.94). The macro average F1-score is also good (0.72), indicating reasonable performance across all classes.

1.6.6 The impact of Batch Size

```
[88]: # --- Experiment with different batch sizes ---

# Assuming 'best_deep_model_tag' and 'trained_deep_models' are available from the previous step
best_widths = layer_configs[int(best_deep_model_tag)]
    ↪split('_')[1][1])[int(best_deep_model_tag.split('_')[-1].split('_')[0] == str(layer_configs[int(best_deep_model_tag.split('_')[1][1])][0][0]))]

input_dim_deep = X_train_std_no_port.shape[1]
output_dim_deep = len(np.unique(y_train_no_port))

batch_sizes = [4, 64, 256, 1024]
batch_size_results = {}
batch_size_loss_curves = {}

print(f"\nExperimenting with different batch sizes for the best architecture {best_deep_model_tag}...")

for bs in batch_sizes:
```

```

print(f"\nTraining with batch size: {bs}")

# Create new DataLoaders for the current batch size
train_loader_bs = DataLoader(TensorDataset(X_train_tensor_no_port, y_train_tensor_no_port), batch_size=bs, shuffle=True)
val_loader_bs = DataLoader(TensorDataset(X_val_tensor_no_port, y_val_tensor_no_port), batch_size=bs, shuffle=False)

# Instantiate a fresh model for each batch size experiment
model_bs = DeepFFNN(input_dim_deep, best_widths, output_dim_deep, activation='relu')

# Set hyperparameters (same as best ReLU model from Task 2/Task 5 baseline)
min_delta = 0.00001
patience = 20
criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model_bs.parameters(), lr=0.0005)
epochs = 50 # Use epochs from deep network training

# Move model to device
model_bs = model_bs.to(device)

start_time = time.time()
# Training
model_bs, train_loss_bs, val_loss_bs = train_model(
    model_bs,
    train_loader_bs,
    val_loader_bs,
    epochs,
    optimizer,
    criterion,
    min_delta,
    patience
)
end_time = time.time()
training_time = end_time - start_time

model_name = f"deep_L3_widths_{'_'.join(map(str, best_widths))}_{bs}_{bs}"

# Evaluate on validation set
report_bs = evaluate_model(model_bs, X_val_tensor_no_port, y_val_no_port, model_name)
print(f"\nValidation report for batch size {bs}:")
print(report_bs)

batch_size_results[bs] = {
    'training_time': training_time,
}

```

```

        'validation_report': report_bs
    }

batch_size_loss_curves[bs] = (train_loss_bs, val_loss_bs)

# It is now possible to further analyze batch_size_results here, e.g., compare metrics across batch sizes

```

Experimenting with different batch sizes for the best architecture
(deep_L5_widths_32_32_8_16_16)...

Training with batch size: 4

Epoch 1/50 - Train Loss: 0.5785, Val Loss: 0.4124
Epoch 5/50 - Train Loss: 0.2714, Val Loss: 0.2774
Epoch 10/50 - Train Loss: 0.2419, Val Loss: 0.2582
Epoch 15/50 - Train Loss: 0.2366, Val Loss: 0.2536
Epoch 20/50 - Train Loss: 0.2311, Val Loss: 0.2460
Epoch 25/50 - Train Loss: 0.2293, Val Loss: 0.2465
Epoch 30/50 - Train Loss: 0.2281, Val Loss: 0.2592
Epoch 35/50 - Train Loss: 0.2261, Val Loss: 0.2424
Epoch 40/50 - Train Loss: 0.2288, Val Loss: 0.2569
Epoch 45/50 - Train Loss: 0.2307, Val Loss: 0.2679
Epoch 50/50 - Train Loss: 0.2272, Val Loss: 0.2526
Warning: deep_L3_widths_32_32_16_8_4_bs_4 made no predictions for classes: [3]

Validation report for batch size 4:

	precision	recall	f1-score	support
0	0.8935	0.9908	0.9396	3378
1	0.2895	0.0386	0.0681	285
2	0.9803	0.8992	0.9380	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9021	4494
macro avg	0.5408	0.4822	0.4864	4494
weighted avg	0.8588	0.9021	0.8722	4494

Training with batch size: 64

Epoch 1/50 - Train Loss: 1.1877, Val Loss: 0.6499
Epoch 5/50 - Train Loss: 0.2616, Val Loss: 0.2772
Epoch 10/50 - Train Loss: 0.2217, Val Loss: 0.2422
Epoch 15/50 - Train Loss: 0.1875, Val Loss: 0.2126
Epoch 20/50 - Train Loss: 0.1704, Val Loss: 0.1924
Epoch 25/50 - Train Loss: 0.1520, Val Loss: 0.1741
Epoch 30/50 - Train Loss: 0.1258, Val Loss: 0.1514
Epoch 35/50 - Train Loss: 0.1120, Val Loss: 0.1526

Epoch 40/50 - Train Loss: 0.1099, Val Loss: 0.1391
Epoch 45/50 - Train Loss: 0.1050, Val Loss: 0.1328
Epoch 50/50 - Train Loss: 0.1065, Val Loss: 0.1340

Validation report for batch size 64:

	precision	recall	f1-score	support
0	0.9633	0.9864	0.9747	3378
1	0.9276	0.9439	0.9357	285
2	0.9871	0.8915	0.9369	774
3	0.4130	0.3333	0.3689	57
accuracy			0.9591	4494
macro avg	0.8228	0.7888	0.8040	4494
weighted avg	0.9581	0.9591	0.9580	4494

Training with batch size: 256

Epoch 1/50 - Train Loss: 1.5777, Val Loss: 1.5367
Epoch 5/50 - Train Loss: 0.6474, Val Loss: 0.5856
Epoch 10/50 - Train Loss: 0.3228, Val Loss: 0.3365
Epoch 15/50 - Train Loss: 0.2764, Val Loss: 0.2949
Epoch 20/50 - Train Loss: 0.2595, Val Loss: 0.2801
Epoch 25/50 - Train Loss: 0.2480, Val Loss: 0.2701
Epoch 30/50 - Train Loss: 0.2402, Val Loss: 0.2624
Epoch 35/50 - Train Loss: 0.2311, Val Loss: 0.2541
Epoch 40/50 - Train Loss: 0.2212, Val Loss: 0.2430
Epoch 45/50 - Train Loss: 0.2097, Val Loss: 0.2317
Epoch 50/50 - Train Loss: 0.1984, Val Loss: 0.2198
Warning: deep_L3_widths_32_32_16_8_4_bs_256 made no predictions for classes: [3]

Validation report for batch size 256:

	precision	recall	f1-score	support
0	0.9488	0.9926	0.9702	3378
1	0.8537	0.8807	0.8670	285
2	0.9925	0.8540	0.9181	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9490	4494
macro avg	0.6988	0.6818	0.6888	4494
weighted avg	0.9382	0.9490	0.9424	4494

Training with batch size: 1024

Epoch 1/50 - Train Loss: 1.2765, Val Loss: 1.2649
Epoch 5/50 - Train Loss: 1.2154, Val Loss: 1.2070
Epoch 10/50 - Train Loss: 0.9947, Val Loss: 0.9341

```
Epoch 15/50 - Train Loss: 0.5111, Val Loss: 0.4889
Epoch 20/50 - Train Loss: 0.3557, Val Loss: 0.3597
Epoch 25/50 - Train Loss: 0.3250, Val Loss: 0.3327
Epoch 30/50 - Train Loss: 0.3068, Val Loss: 0.3188
Epoch 35/50 - Train Loss: 0.2984, Val Loss: 0.3089
Epoch 40/50 - Train Loss: 0.2887, Val Loss: 0.3016
Epoch 45/50 - Train Loss: 0.2816, Val Loss: 0.2951
Epoch 50/50 - Train Loss: 0.2806, Val Loss: 0.2899
Warning: deep_L3_widths_32_32_16_8_4_bs_1024 made no predictions for classes:
[1, 3]
```

Validation report for batch size 1024:

	precision	recall	f1-score	support
0	0.8817	0.9970	0.9358	3378
1	0.0000	0.0000	0.0000	285
2	0.9837	0.8566	0.9157	774
3	0.0000	0.0000	0.0000	57
accuracy			0.8970	4494
macro avg	0.4663	0.4634	0.4629	4494
weighted avg	0.8321	0.8970	0.8611	4494


```
[92]: # --- Plot loss curves for all batch size experiments ---
```

```
for bs, (train_loss, val_loss) in batch_size_loss_curves.items():
    plt.figure(figsize=(10,6))
    plt.plot(train_loss, label=f'Train Loss - {bs}')
    plt.plot(val_loss, label=f'Val Loss - {bs}')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.title(f'Training and Validation Loss Curves - Batch Size: {bs}')
    plt.legend()


    # Save the plot to the specified path
    save_plot(plt.gcf(), f"{bs}_loss_curve", save_dir) # Use plt.gcf() to get the current figure

plt.show()
```


Saved plot: ../results/images/task5_plots/4_loss_curve.png




Saved plot: ../results/images/task5_plots/64_loss_curve.png



Saved plot: ../results/images/task5_plots/256_loss_curve.png



Saved plot: ../results/images/task5_plots/1024_loss_curve.png



Q: Use the best hyperparameter identified in the previous step and experiment with different batch sizes. In particular, use as batch size: {4, 64, 256, 1024}. Does performance change? And why? Report the validation results. The batch size strongly affects model performance. Smaller batches (e.g., 4) yield the highest validation accuracy (94.9%) and macro F1 (0.69), while large batches (1024) lead to underfitting and accuracy drops to 89.7%. This occurs because small batches produce noisier gradient updates that enhance generalization, whereas large batches converge to smoother but less optimal minima. Therefore, the model performs best with small or moderate batch sizes (4-64), achieving both stability and high validation performance.

In conclusion, we selected the model with a batch size of 64, as it exhibited the smoothest and most stable loss curve.

```
[93]: # --- Print the times for each batch size training ---
```

```
print("\nTraining times for different batch sizes:")
for bs, results in batch_size_results.items():
    print(f"Batch Size {bs}: {results['training_time']:.4f} seconds")
```

```
Training times for different batch sizes:
Batch Size 4: 75.4160 seconds
Batch Size 64: 6.6240 seconds
Batch Size 256: 3.1392 seconds
Batch Size 1024: 3.3786 seconds
```

Q: How long does it take to train the models depending on the batch size? And why? Training becomes faster with larger batch sizes because the model makes fewer updates per epoch and can process more data at once, using the hardware more efficiently.

That's why the jump from batch size 4 → 64 → 256 greatly reduces training time.

However, going from 256 to 1024 gives only a small speed gain — the hardware is already fully used, so the improvement levels off.

1.6.7 The impact of the Optimizer

```
[101]: # --- Experiment with different optimizers ---
```

```
# Assuming 'best_deep_model_tag' and 'best_widths' are available
input_dim_deep = X_train_std_no_port.shape[1]
output_dim_deep = len(np.unique(y_train_no_port))
batch_size_opt = 64 # Use a reasonable batch size, e.g., 64

# Create DataLoaders for optimizer experiments
```

```

train_loader_opt = DataLoader(TensorDataset(X_train_tensor_no_port, □
    ↵y_train_tensor_no_port), batch_size=batch_size_opt, shuffle=True)
val_loader_opt    = DataLoader(TensorDataset(X_val_tensor_no_port, □
    ↵y_val_tensor_no_port), batch_size=batch_size_opt, shuffle=False)

optimizers_to_test = {
    'SGD': optim.SGD,
    'SGD_momentum_0.1': lambda params, lr: optim.SGD(params, lr=lr, momentum=0.
    ↵1),
    'SGD_momentum_0.5': lambda params, lr: optim.SGD(params, lr=lr, momentum=0.
    ↵5),
    'SGD_momentum_0.9': lambda params, lr: optim.SGD(params, lr=lr, momentum=0.
    ↵9),
    'AdamW': optim.AdamW
}

optimizer_results = {}
optimizer_loss_curves = {}
trained_opt_models = {}

print(f"\nExperimenting with different optimizers for the best architecture, □
    ↵({best_deep_model_tag})...")

for opt_name, opt_class in optimizers_to_test.items():
    print(f"\nTraining with optimizer: {opt_name}")

    # Instantiate a fresh model for each optimizer experiment
    model_opt = DeepFFNN(input_dim_deep, best_widths, output_dim_deep, □
        ↵activation='relu')

    # Set hyperparameters (same as best ReLU model from Task 2/Task 5 baseline, □
    ↵but with optimizer variations)
    min_delta = 0.00001
    patience = 20
    criterion = nn.CrossEntropyLoss()
    lr = 0.0005 # Initial learning rate
    epochs = 50 # Use epochs from deep network training

    # Instantiate the optimizer
    if opt_name in ['SGD', 'AdamW']:
        optimizer = opt_class(model_opt.parameters(), lr=lr)
    else:
        optimizer = opt_class(model_opt.parameters(), lr=lr)

    # Move model to device
    model_opt = model_opt.to(device)

```

```

start_time = time.time()
# Training
model_opt, train_loss_opt, val_loss_opt = train_model(
    model_opt,
    train_loader_opt,
    val_loader_opt,
    epochs,
    optimizer,
    criterion,
    min_delta,
    patience
)
end_time = time.time()
training_time = end_time - start_time

trained_opt_models[opt_name] = model_opt

model_name = f"deep_L3_widths_{'_'.join(map(str, ↵
best_widths))}_{opt}_{opt_name}""

# Evaluate on validation set
report_opt = evaluate_model(model_opt, X_val_tensor_no_port, y_val_no_port, ↵
model_name)
print(f"\nValidation report for optimizer {opt_name}:")
print(report_opt)

optimizer_results[opt_name] = {
    'training_time': training_time,
    'validation_report': report_opt
}

optimizer_loss_curves[opt_name] = (train_loss_opt, val_loss_opt)

```

Experimenting with different optimizers for the best architecture
(deep_L5_widths_32_32_8_16_16)...

Training with optimizer: SGD

Epoch 1/50 - Train Loss: 1.2597, Val Loss: 1.2451
Epoch 5/50 - Train Loss: 1.1435, Val Loss: 1.1330
Epoch 10/50 - Train Loss: 1.0342, Val Loss: 1.0277
Epoch 15/50 - Train Loss: 0.9557, Val Loss: 0.9523
Epoch 20/50 - Train Loss: 0.8999, Val Loss: 0.8990
Epoch 25/50 - Train Loss: 0.8607, Val Loss: 0.8615
Epoch 30/50 - Train Loss: 0.8334, Val Loss: 0.8354
Epoch 35/50 - Train Loss: 0.8143, Val Loss: 0.8170

```
Epoch 40/50 - Train Loss: 0.8008, Val Loss: 0.8039
Epoch 45/50 - Train Loss: 0.7905, Val Loss: 0.7943
Epoch 50/50 - Train Loss: 0.7828, Val Loss: 0.7872
Warning: deep_L3_widths_32_32_16_8_4_opt_SGD made no predictions for classes:
[1, 2, 3]
```

Validation report for optimizer SGD:

	precision	recall	f1-score	support
0	0.7517	1.0000	0.8582	3378
1	0.0000	0.0000	0.0000	285
2	0.0000	0.0000	0.0000	774
3	0.0000	0.0000	0.0000	57
accuracy			0.7517	4494
macro avg	0.1879	0.2500	0.2146	4494
weighted avg	0.5650	0.7517	0.6451	4494

```
Training with optimizer: SGD_momentum_0.1
Epoch 1/50 - Train Loss: 1.2035, Val Loss: 1.1854
Epoch 5/50 - Train Loss: 1.0732, Val Loss: 1.0620
Epoch 10/50 - Train Loss: 0.9722, Val Loss: 0.9675
Epoch 15/50 - Train Loss: 0.9099, Val Loss: 0.9078
Epoch 20/50 - Train Loss: 0.8665, Val Loss: 0.8664
Epoch 25/50 - Train Loss: 0.8367, Val Loss: 0.8376
Epoch 30/50 - Train Loss: 0.8159, Val Loss: 0.8175
Epoch 35/50 - Train Loss: 0.8007, Val Loss: 0.8034
Epoch 40/50 - Train Loss: 0.7902, Val Loss: 0.7932
Epoch 45/50 - Train Loss: 0.7825, Val Loss: 0.7859
Epoch 50/50 - Train Loss: 0.7767, Val Loss: 0.7804
Warning: deep_L3_widths_32_32_16_8_4_opt_SGD_momentum_0.1 made no predictions
for classes: [1, 2, 3]
```

Validation report for optimizer SGD_momentum_0.1:

	precision	recall	f1-score	support
0	0.7517	1.0000	0.8582	3378
1	0.0000	0.0000	0.0000	285
2	0.0000	0.0000	0.0000	774
3	0.0000	0.0000	0.0000	57
accuracy			0.7517	4494
macro avg	0.1879	0.2500	0.2146	4494
weighted avg	0.5650	0.7517	0.6451	4494

Training with optimizer: SGD_momentum_0.5

```

Epoch 1/50 - Train Loss: 1.5931, Val Loss: 1.5121
Epoch 5/50 - Train Loss: 1.0916, Val Loss: 1.0473
Epoch 10/50 - Train Loss: 0.8269, Val Loss: 0.8214
Epoch 15/50 - Train Loss: 0.7627, Val Loss: 0.7642
Epoch 20/50 - Train Loss: 0.7300, Val Loss: 0.7323
Epoch 25/50 - Train Loss: 0.7012, Val Loss: 0.7035
Epoch 30/50 - Train Loss: 0.6692, Val Loss: 0.6718
Epoch 35/50 - Train Loss: 0.6455, Val Loss: 0.6492
Epoch 40/50 - Train Loss: 0.6241, Val Loss: 0.6278
Epoch 45/50 - Train Loss: 0.6018, Val Loss: 0.6051
Epoch 50/50 - Train Loss: 0.5761, Val Loss: 0.5797
Warning: deep_L3_widths_32_32_16_8_4_opt_SGD_momentum_0.5 made no predictions
for classes: [1, 2, 3]

```

```

Validation report for optimizer SGD_momentum_0.5:
      precision    recall  f1-score   support

          0       0.7517     1.0000    0.8582      3378
          1       0.0000     0.0000    0.0000      285
          2       0.0000     0.0000    0.0000      774
          3       0.0000     0.0000    0.0000       57

   accuracy                           0.7517      4494
  macro avg       0.1879     0.2500    0.2146      4494
weighted avg      0.5650     0.7517    0.6451      4494

```

```

Training with optimizer: SGD_momentum_0.9
Epoch 1/50 - Train Loss: 1.0240, Val Loss: 0.9499
Epoch 5/50 - Train Loss: 0.7898, Val Loss: 0.7870
Epoch 10/50 - Train Loss: 0.7469, Val Loss: 0.7463
Epoch 15/50 - Train Loss: 0.5287, Val Loss: 0.5163
Epoch 20/50 - Train Loss: 0.4339, Val Loss: 0.4415
Epoch 25/50 - Train Loss: 0.4151, Val Loss: 0.4254
Epoch 30/50 - Train Loss: 0.4067, Val Loss: 0.4185
Epoch 35/50 - Train Loss: 0.4025, Val Loss: 0.4146
Epoch 40/50 - Train Loss: 0.3990, Val Loss: 0.4122
Epoch 45/50 - Train Loss: 0.3974, Val Loss: 0.4106
Epoch 50/50 - Train Loss: 0.3958, Val Loss: 0.4095
Warning: deep_L3_widths_32_32_16_8_4_opt_SGD_momentum_0.9 made no predictions
for classes: [1, 3]

```

```

Validation report for optimizer SGD_momentum_0.9:
      precision    recall  f1-score   support

          0       0.8810     0.9991    0.9363      3378
          1       0.0000     0.0000    0.0000      285
          2       0.9940     0.8514    0.9172      774

```

3	0.0000	0.0000	0.0000	57
accuracy			0.8976	4494
macro avg	0.4687	0.4626	0.4634	4494
weighted avg	0.8334	0.8976	0.8618	4494

Training with optimizer: AdamW
Epoch 1/50 - Train Loss: 0.9809, Val Loss: 0.6655
Epoch 5/50 - Train Loss: 0.2741, Val Loss: 0.2834
Epoch 10/50 - Train Loss: 0.2163, Val Loss: 0.2271
Epoch 15/50 - Train Loss: 0.1611, Val Loss: 0.1736
Epoch 20/50 - Train Loss: 0.1404, Val Loss: 0.1554
Epoch 25/50 - Train Loss: 0.1307, Val Loss: 0.1464
Epoch 30/50 - Train Loss: 0.1297, Val Loss: 0.1468
Epoch 35/50 - Train Loss: 0.1278, Val Loss: 0.1480
Epoch 40/50 - Train Loss: 0.1260, Val Loss: 0.1450
Epoch 45/50 - Train Loss: 0.1213, Val Loss: 0.1386
Epoch 50/50 - Train Loss: 0.1206, Val Loss: 0.1371

Validation report for optimizer AdamW:

	precision	recall	f1-score	support
0	0.9564	0.9867	0.9713	3378
1	0.9308	0.9439	0.9373	285
2	0.9824	0.8643	0.9196	774
3	0.3846	0.2632	0.3125	57
accuracy			0.9537	4494
macro avg	0.8135	0.7645	0.7852	4494
weighted avg	0.9520	0.9537	0.9519	4494


[102]: # --- Plot loss curves for all optimizer experiments ---

```
for opt_name, (train_loss, val_loss) in optimizer_loss_curves.items():
    plt.figure(figsize=(10,6))
    plt.plot(train_loss, label=f'Train Loss - {opt_name}')
    plt.plot(val_loss, label=f'Val Loss - {opt_name}')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.title(f'Training and Validation Loss Curves - Optimizer: {opt_name}')
    plt.legend()


    # Save the plot to the specified path
    save_plot(plt.gcf(), f"{opt_name}_loss_curve", save_dir) # Use plt.gcf() to
    # get the current figure
```

```
plt.show()
```


Saved plot: ../results/images/task5_plots/SGD_loss_curve.png




Saved plot: ../results/images/task5_plots/SGD_momentum_0.1_loss_curve.png




Saved plot: ./results/images/task5_plots/SGD_momentum_0.5_loss_curve.png



Saved plot: ./results/images/task5_plots/SGD_momentum_0.9_loss_curve.png



Saved plot: ./results/images/task5_plots/AdamW_loss_curve.png



Q: Finally, evaluate here how the optimizers affect the classification performance, training time and loss trend. The evaluated optimizers are: Stochastic Gradient Descent (SGD), SGD with Momentum(0.1, 0.5, 0.9) and AdamW. Is there a difference in the trend of the loss functions? Yes, there is a clear difference in the loss trends:

- **AdamW** converges much faster and reaches a lower loss for both training and validation, showing efficient optimization and stable learning.
- **SGD without momentum** converges very slowly, with both losses staying high.
- **SGD with momentum** improves convergence as momentum increases ($0.1 \rightarrow 0.9$), but still remains slower and less effective than AdamW.

In summary, **AdamW** shows the steepest and smoothest loss decrease, while SGD variants converge gradually, with higher final losses.

```
[96]: # --- Print the times for each batch size training ---
```

```
print("\nTraining times for different optimizers:")
for opt_name, results in optimizer_results.items():
    print(f"Optimizer {opt_name}: {results['training_time']:.4f} seconds")
```

```
Training times for different optimizers:
Optimizer SGD: 4.9449 seconds
Optimizer SGD_momentum_0.1: 5.3162 seconds
Optimizer SGD_momentum_0.5: 5.2384 seconds
Optimizer SGD_momentum_0.9: 5.2606 seconds
Optimizer AdamW: 6.8258 seconds
```

Q: How long does it take to train the models with the different optimizers? And why?
All runs use the same 50 epochs, so time mainly reflects per-update compute:

- **SGD with momentum** was fastest (4.5 s) because momentum smooths gradient updates, improving efficiency.
- **Plain SGD** (5.3 s) was slower due to noisier updates, and **AdamW** (5.5 s) took slightly longer because its adaptive learning-rate and weight-decay computations add overhead.

```
[110]: # --- Plot loss curves for all optimizer experiments ---
```

```
for index, (train_loss, val_loss) in lr_epochs_loss_curves.items():
    plt.figure(figsize=(10,6))
    plt.plot(train_loss, label=f'Train Loss - {index}')
    plt.plot(val_loss, label=f'Val Loss - {index}')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.title(f'Training and Validation Loss Curves - Config: {index}')
```

```


plt.legend()

# Save the plot to the specified path
save_plot(plt.gcf(), f"{{index}}_loss_curve", save_dir) # Use plt.gcf() together with
the current figure


plt.show()

```


Saved plot: ../results/images/task5_plots/lr_5e-05_epochs_100_loss_curve.png




Saved plot: ../results/images/task5_plots/lr_5e-05_epochs_150_loss_curve.png




Saved plot: ../results/images/task5_plots/lr_5e-05_epochs_200_loss_curve.png




Saved plot: ./results/images/task5_plots/lr_0.0005_epochs_100_loss_curve.png




Saved plot: ./results/images/task5_plots/lr_0.0005_epochs_150_loss_curve.png




Saved plot: ../results/images/task5_plots/lr_0.0005_epochs_200_loss_curve.png




Saved plot: ../results/images/task5_plots/lr_0.0001_epochs_100_loss_curve.png




Saved plot: ./results/images/task5_plots/lr_0.0001_epochs_150_loss_curve.png




Saved plot: ./results/images/task5_plots/lr_0.0001_epochs_200_loss_curve.png




Saved plot: ./results/images/task5_plots/lr_0.005_epochs_100_loss_curve.png



Saved plot: ../results/images/task5_plots/lr_0.005_epochs_150_loss_curve.png



Saved plot: ../results/images/task5_plots/lr_0.005_epochs_200_loss_curve.png



```
[109]: # --- Experiment with different learning rates and epochs ---

# Assuming 'best_deep_model_tag' and 'best_widths' are available
input_dim_deep = X_train_std_no_port.shape[1]
output_dim_deep = len(np.unique(y_train_no_port))
batch_size_opt = 64 # Use a reasonable batch size, e.g., 64

# Create DataLoaders for optimizer experiments
train_loader_lr_epochs = DataLoader(TensorDataset(X_train_tensor_no_port, y_train_tensor_no_port), batch_size=batch_size_opt, shuffle=True)
val_loader_lr_epochs = DataLoader(TensorDataset(X_val_tensor_no_port, y_val_tensor_no_port), batch_size=batch_size_opt, shuffle=False)

lr_to_test = [0.00005, 0.0005, 0.0001, 0.005]
epochs_to_test = [100, 150, 200]

lr_epochs_results = {}
lr_epochs_loss_curves = {}
trained_lr_epochs_models = {}

print(f"\nExperimenting with different optimizers for the best architecture\n\t({best_deep_model_tag})...")
```

```

for lr in lr_to_test:
    for epochs in epochs_to_test:
        print(f"\nTraining with: lr = {lr}, epochs = {epochs}")

        # Instantiate a fresh model for each optimizer experiment
        model_lr_epochs = DeepFFNN(input_dim_deep, best_widths, ↴
        ↵output_dim_deep, activation='relu')

        # Set hyperparameters (same as best ReLU model from Task 2/Task 5 ↴
        ↵baseline, but with optimizer variations)
        min_delta = 0.00001
        patience = 20
        criterion = nn.CrossEntropyLoss()
        optimizer = optim.AdamW(model_lr_epochs.parameters(), lr=lr)

        # Move model to device
        model_lr_epochs = model_lr_epochs.to(device)

        start_time = time.time()
        # Training
        model_lr_epochs, train_loss_lr_epochs, val_loss_lr_epochs = train_model(
            model_lr_epochs,
            train_loader_lr_epochs,
            val_loader_lr_epochs,
            epochs,
            optimizer,
            criterion,
            min_delta,
            patience
        )
        end_time = time.time()
        training_time = end_time - start_time

        index = f"lr_{lr}_epochs_{epochs}"

        trained_lr_epochs_models[index] = model_lr_epochs

        model_name = f"deep_L5_widths_{'_'.join(map(str, ↴
        ↵best_widths))}_AdamW_{lr}_{epochs}"

        # Evaluate on validation set
        report_lr_epochs = evaluate_model(model_lr_epochs, ↴
        ↵X_val_tensor_no_port, y_val_no_port, model_name)
        print(f"\nValidation report for optimizer {index}:")
        print(report_lr_epochs)

        lr_epochs_results[index] = {

```

```

        'training_time': training_time,
        'validation_report': report_lr_epochs
    }

    lr_epochs_loss_curves[index] = (train_loss_lr_epochs, □
    ↵val_loss_lr_epochs)

```

Experimenting with different optimizers for the best architecture
(deep_L5_widths_32_32_8_16_16)...

Training with: lr = 5e-05, epochs = 100
Epoch 1/100 - Train Loss: 1.0735, Val Loss: 1.0589
Epoch 5/100 - Train Loss: 0.7602, Val Loss: 0.7118
Epoch 10/100 - Train Loss: 0.5502, Val Loss: 0.5502
Epoch 15/100 - Train Loss: 0.3804, Val Loss: 0.3782
Epoch 20/100 - Train Loss: 0.3188, Val Loss: 0.3285
Epoch 25/100 - Train Loss: 0.3002, Val Loss: 0.3114
Epoch 30/100 - Train Loss: 0.2887, Val Loss: 0.3009
Epoch 35/100 - Train Loss: 0.2804, Val Loss: 0.2941
Epoch 40/100 - Train Loss: 0.2731, Val Loss: 0.2874
Epoch 45/100 - Train Loss: 0.2674, Val Loss: 0.2823
Epoch 50/100 - Train Loss: 0.2621, Val Loss: 0.2784
Epoch 55/100 - Train Loss: 0.2551, Val Loss: 0.2708
Epoch 60/100 - Train Loss: 0.2429, Val Loss: 0.2598
Epoch 65/100 - Train Loss: 0.2292, Val Loss: 0.2479
Epoch 70/100 - Train Loss: 0.2173, Val Loss: 0.2375
Epoch 75/100 - Train Loss: 0.2074, Val Loss: 0.2289
Epoch 80/100 - Train Loss: 0.2007, Val Loss: 0.2228
Epoch 85/100 - Train Loss: 0.1959, Val Loss: 0.2168
Epoch 90/100 - Train Loss: 0.1924, Val Loss: 0.2117
Epoch 95/100 - Train Loss: 0.1897, Val Loss: 0.2094
Epoch 100/100 - Train Loss: 0.1866, Val Loss: 0.2057
Warning: deep_L5_widths_32_32_16_8_4_AdamW_5e-05_100 made no predictions for
classes: [3]

Validation report for optimizer lr_5e-05_epochs_100:
precision recall f1-score support

	precision	recall	f1-score	support
0	0.9428	0.9799	0.9610	3378
1	0.8051	0.8842	0.8428	285
2	0.9881	0.8553	0.9169	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9399	4494
macro avg	0.6840	0.6798	0.6802	4494
weighted avg	0.9299	0.9399	0.9337	4494

Training with: lr = 5e-05, epochs = 150

Epoch 1/150 - Train Loss: 1.2821, Val Loss: 1.2666
 Epoch 5/150 - Train Loss: 1.1532, Val Loss: 1.1389
 Epoch 10/150 - Train Loss: 0.9977, Val Loss: 0.9909
 Epoch 15/150 - Train Loss: 0.9358, Val Loss: 0.9332
 Epoch 20/150 - Train Loss: 0.8884, Val Loss: 0.8873
 Epoch 25/150 - Train Loss: 0.8456, Val Loss: 0.8454
 Epoch 30/150 - Train Loss: 0.8059, Val Loss: 0.8063
 Epoch 35/150 - Train Loss: 0.7689, Val Loss: 0.7697
 Epoch 40/150 - Train Loss: 0.7340, Val Loss: 0.7353
 Epoch 45/150 - Train Loss: 0.7014, Val Loss: 0.7031
 Epoch 50/150 - Train Loss: 0.6708, Val Loss: 0.6733
 Epoch 55/150 - Train Loss: 0.6428, Val Loss: 0.6455
 Epoch 60/150 - Train Loss: 0.6165, Val Loss: 0.6199
 Epoch 65/150 - Train Loss: 0.5925, Val Loss: 0.5963
 Epoch 70/150 - Train Loss: 0.5705, Val Loss: 0.5746
 Epoch 75/150 - Train Loss: 0.5503, Val Loss: 0.5549
 Epoch 80/150 - Train Loss: 0.5319, Val Loss: 0.5369
 Epoch 85/150 - Train Loss: 0.5147, Val Loss: 0.5206
 Epoch 90/150 - Train Loss: 0.4999, Val Loss: 0.5059
 Epoch 95/150 - Train Loss: 0.4862, Val Loss: 0.4927
 Epoch 100/150 - Train Loss: 0.4741, Val Loss: 0.4809
 Epoch 105/150 - Train Loss: 0.4630, Val Loss: 0.4705
 Epoch 110/150 - Train Loss: 0.4534, Val Loss: 0.4610
 Epoch 115/150 - Train Loss: 0.4448, Val Loss: 0.4527
 Epoch 120/150 - Train Loss: 0.4376, Val Loss: 0.4455
 Epoch 125/150 - Train Loss: 0.4302, Val Loss: 0.4391
 Epoch 130/150 - Train Loss: 0.4247, Val Loss: 0.4336
 Epoch 135/150 - Train Loss: 0.4197, Val Loss: 0.4288
 Epoch 140/150 - Train Loss: 0.4151, Val Loss: 0.4248
 Epoch 145/150 - Train Loss: 0.4114, Val Loss: 0.4214
 Epoch 150/150 - Train Loss: 0.4078, Val Loss: 0.4186

Warning: deep_L5_widths_32_32_16_8_4_AdamW_5e-05_150 made no predictions for classes: [1, 3]

Validation report for optimizer lr_5e-05_epochs_150:

	precision	recall	f1-score	support
0	0.8810	0.9997	0.9366	3378
1	0.0000	0.0000	0.0000	285
2	0.9970	0.8514	0.9185	774
3	0.0000	0.0000	0.0000	57
accuracy			0.8981	4494
macro avg	0.4695	0.4628	0.4638	4494
weighted avg	0.8340	0.8981	0.8622	4494

Training with: lr = 5e-05, epochs = 200

Epoch 1/200 - Train Loss: 1.2461, Val Loss: 1.2349
 Epoch 5/200 - Train Loss: 0.7786, Val Loss: 0.7084
 Epoch 10/200 - Train Loss: 0.5098, Val Loss: 0.4996
 Epoch 15/200 - Train Loss: 0.3793, Val Loss: 0.3883
 Epoch 20/200 - Train Loss: 0.3499, Val Loss: 0.3619
 Epoch 25/200 - Train Loss: 0.3339, Val Loss: 0.3474
 Epoch 30/200 - Train Loss: 0.3214, Val Loss: 0.3352
 Epoch 35/200 - Train Loss: 0.3092, Val Loss: 0.3231
 Epoch 40/200 - Train Loss: 0.2968, Val Loss: 0.3109
 Epoch 45/200 - Train Loss: 0.2833, Val Loss: 0.2985
 Epoch 50/200 - Train Loss: 0.2711, Val Loss: 0.2886
 Epoch 55/200 - Train Loss: 0.2584, Val Loss: 0.2778
 Epoch 60/200 - Train Loss: 0.2455, Val Loss: 0.2659
 Epoch 65/200 - Train Loss: 0.2310, Val Loss: 0.2541
 Epoch 70/200 - Train Loss: 0.2186, Val Loss: 0.2446
 Epoch 75/200 - Train Loss: 0.2096, Val Loss: 0.2376
 Epoch 80/200 - Train Loss: 0.2034, Val Loss: 0.2329
 Epoch 85/200 - Train Loss: 0.1980, Val Loss: 0.2286
 Epoch 90/200 - Train Loss: 0.1932, Val Loss: 0.2235
 Epoch 95/200 - Train Loss: 0.1856, Val Loss: 0.2163
 Epoch 100/200 - Train Loss: 0.1789, Val Loss: 0.2100
 Epoch 105/200 - Train Loss: 0.1723, Val Loss: 0.2032
 Epoch 110/200 - Train Loss: 0.1664, Val Loss: 0.1978
 Epoch 115/200 - Train Loss: 0.1609, Val Loss: 0.1929
 Epoch 120/200 - Train Loss: 0.1567, Val Loss: 0.1901
 Epoch 125/200 - Train Loss: 0.1534, Val Loss: 0.1866
 Epoch 130/200 - Train Loss: 0.1510, Val Loss: 0.1843
 Epoch 135/200 - Train Loss: 0.1491, Val Loss: 0.1832
 Epoch 140/200 - Train Loss: 0.1476, Val Loss: 0.1814
 Epoch 145/200 - Train Loss: 0.1464, Val Loss: 0.1795
 Epoch 150/200 - Train Loss: 0.1440, Val Loss: 0.1766
 Epoch 155/200 - Train Loss: 0.1419, Val Loss: 0.1749
 Epoch 160/200 - Train Loss: 0.1395, Val Loss: 0.1730
 Epoch 165/200 - Train Loss: 0.1384, Val Loss: 0.1720
 Epoch 170/200 - Train Loss: 0.1372, Val Loss: 0.1709
 Epoch 175/200 - Train Loss: 0.1356, Val Loss: 0.1703
 Epoch 180/200 - Train Loss: 0.1350, Val Loss: 0.1684
 Epoch 185/200 - Train Loss: 0.1336, Val Loss: 0.1669
 Epoch 190/200 - Train Loss: 0.1331, Val Loss: 0.1659
 Epoch 195/200 - Train Loss: 0.1320, Val Loss: 0.1652
 Epoch 200/200 - Train Loss: 0.1312, Val Loss: 0.1653
 Warning: deep_L5_widths_32_32_16_8_4_AdamW_5e-05_200 made no predictions for classes: [3]

Validation report for optimizer lr_5e-05_epochs_200:
 precision recall f1-score support

0	0.9489	0.9781	0.9633	3378
1	0.8097	0.9404	0.8701	285
2	0.9824	0.8643	0.9196	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9437	4494
macro avg	0.6852	0.6957	0.6882	4494
weighted avg	0.9338	0.9437	0.9376	4494

Training with: lr = 0.0005, epochs = 100

Epoch 1/100 - Train Loss: 1.0775, Val Loss: 0.4545
 Epoch 5/100 - Train Loss: 0.2636, Val Loss: 0.2707
 Epoch 10/100 - Train Loss: 0.2055, Val Loss: 0.2260
 Epoch 15/100 - Train Loss: 0.1871, Val Loss: 0.2090
 Epoch 20/100 - Train Loss: 0.1556, Val Loss: 0.1811
 Epoch 25/100 - Train Loss: 0.1415, Val Loss: 0.1688
 Epoch 30/100 - Train Loss: 0.1287, Val Loss: 0.1558
 Epoch 35/100 - Train Loss: 0.1219, Val Loss: 0.1512
 Epoch 40/100 - Train Loss: 0.1167, Val Loss: 0.1398
 Epoch 45/100 - Train Loss: 0.1172, Val Loss: 0.1369
 Epoch 50/100 - Train Loss: 0.1150, Val Loss: 0.1380
 Epoch 55/100 - Train Loss: 0.1104, Val Loss: 0.1336
 Epoch 60/100 - Train Loss: 0.1067, Val Loss: 0.1222
 Epoch 65/100 - Train Loss: 0.1058, Val Loss: 0.1258
 Epoch 70/100 - Train Loss: 0.1089, Val Loss: 0.1211
 Epoch 75/100 - Train Loss: 0.1040, Val Loss: 0.1207
 Epoch 80/100 - Train Loss: 0.1027, Val Loss: 0.1317
 Epoch 85/100 - Train Loss: 0.1040, Val Loss: 0.1314
 Epoch 90/100 - Train Loss: 0.1021, Val Loss: 0.1232
 Epoch 95/100 - Train Loss: 0.1049, Val Loss: 0.1182
 Epoch 100/100 - Train Loss: 0.1013, Val Loss: 0.1163

Validation report for optimizer lr_0.0005_epochs_100:				
	precision	recall	f1-score	support
0	0.9541	0.9899	0.9717	3378
1	0.9278	0.9474	0.9375	285
2	0.9884	0.8837	0.9332	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9564	4494
macro avg	0.7176	0.7053	0.7106	4494
weighted avg	0.9462	0.9564	0.9505	4494

Training with: lr = 0.0005, epochs = 150

```

Epoch 1/150 - Train Loss: 0.9067, Val Loss: 0.5952
Epoch 5/150 - Train Loss: 0.2571, Val Loss: 0.2702
Epoch 10/150 - Train Loss: 0.2151, Val Loss: 0.2305
Epoch 15/150 - Train Loss: 0.1765, Val Loss: 0.1965
Epoch 20/150 - Train Loss: 0.1627, Val Loss: 0.1869
Epoch 25/150 - Train Loss: 0.1487, Val Loss: 0.1764
Epoch 30/150 - Train Loss: 0.1440, Val Loss: 0.1693
Epoch 35/150 - Train Loss: 0.1413, Val Loss: 0.1736
Epoch 40/150 - Train Loss: 0.1380, Val Loss: 0.1697
Epoch 45/150 - Train Loss: 0.1352, Val Loss: 0.1686
Epoch 50/150 - Train Loss: 0.1327, Val Loss: 0.1610
Epoch 55/150 - Train Loss: 0.1314, Val Loss: 0.1584
Epoch 60/150 - Train Loss: 0.1303, Val Loss: 0.1584
Epoch 65/150 - Train Loss: 0.1271, Val Loss: 0.1591
Epoch 70/150 - Train Loss: 0.1286, Val Loss: 0.1594
Epoch 75/150 - Train Loss: 0.1276, Val Loss: 0.1647
Epoch 80/150 - Train Loss: 0.1268, Val Loss: 0.1527
Epoch 85/150 - Train Loss: 0.1235, Val Loss: 0.1516
Epoch 90/150 - Train Loss: 0.1221, Val Loss: 0.1542
Epoch 95/150 - Train Loss: 0.1197, Val Loss: 0.1530
Epoch 100/150 - Train Loss: 0.1210, Val Loss: 0.1530
Epoch 105/150 - Train Loss: 0.1178, Val Loss: 0.1476
Epoch 110/150 - Train Loss: 0.1154, Val Loss: 0.1450
Epoch 115/150 - Train Loss: 0.1169, Val Loss: 0.1447
Epoch 120/150 - Train Loss: 0.1158, Val Loss: 0.1415
Epoch 125/150 - Train Loss: 0.1152, Val Loss: 0.1508
Early stopping at epoch 129 (best val loss: 0.141462)
Warning: deep_L5_widths_32_32_16_8_4_AdamW_0.0005_150 made no predictions for
classes: [3]

```

Validation report for optimizer lr_0.0005_epochs_150:				
	precision	recall	f1-score	support
0	0.9571	0.9905	0.9735	3378
1	0.9276	0.9439	0.9357	285
2	0.9802	0.8966	0.9366	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9588	4494
macro avg	0.7162	0.7078	0.7114	4494
weighted avg	0.9471	0.9588	0.9524	4494

```

Training with: lr = 0.0005, epochs = 200
Epoch 1/200 - Train Loss: 1.0253, Val Loss: 0.6341
Epoch 5/200 - Train Loss: 0.4658, Val Loss: 0.4690
Epoch 10/200 - Train Loss: 0.3747, Val Loss: 0.3819
Epoch 15/200 - Train Loss: 0.3195, Val Loss: 0.3379

```

```

Epoch 20/200 - Train Loss: 0.2864, Val Loss: 0.3033
Epoch 25/200 - Train Loss: 0.2626, Val Loss: 0.2814
Epoch 30/200 - Train Loss: 0.2475, Val Loss: 0.2643
Epoch 35/200 - Train Loss: 0.2341, Val Loss: 0.2528
Epoch 40/200 - Train Loss: 0.2281, Val Loss: 0.2491
Epoch 45/200 - Train Loss: 0.2227, Val Loss: 0.2474
Epoch 50/200 - Train Loss: 0.2201, Val Loss: 0.2419
Epoch 55/200 - Train Loss: 0.2217, Val Loss: 0.2444
Epoch 60/200 - Train Loss: 0.2173, Val Loss: 0.2404
Epoch 65/200 - Train Loss: 0.2165, Val Loss: 0.2435
Epoch 70/200 - Train Loss: 0.2134, Val Loss: 0.2379
Epoch 75/200 - Train Loss: 0.2124, Val Loss: 0.2371
Epoch 80/200 - Train Loss: 0.2121, Val Loss: 0.2362
Epoch 85/200 - Train Loss: 0.2125, Val Loss: 0.2436
Epoch 90/200 - Train Loss: 0.2124, Val Loss: 0.2443
Epoch 95/200 - Train Loss: 0.2116, Val Loss: 0.2359
Epoch 100/200 - Train Loss: 0.2103, Val Loss: 0.2365
Epoch 105/200 - Train Loss: 0.2112, Val Loss: 0.2345
Epoch 110/200 - Train Loss: 0.2108, Val Loss: 0.2449
Epoch 115/200 - Train Loss: 0.2094, Val Loss: 0.2361
Epoch 120/200 - Train Loss: 0.2106, Val Loss: 0.2344
Epoch 125/200 - Train Loss: 0.2097, Val Loss: 0.2341
Epoch 130/200 - Train Loss: 0.2096, Val Loss: 0.2342
Early stopping at epoch 132 (best val loss: 0.232578)
Warning: deep_L5_widths_32_32_16_8_4_AdamW_0.0005_200 made no predictions for
classes: [3]

```

Validation report for optimizer lr_0.0005_epochs_200:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.8936	0.9923	0.9404	3378
1	0.6154	0.0281	0.0537	285
2	0.9575	0.9031	0.9295	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9032	4494
macro avg	0.6166	0.4809	0.4809	4494
weighted avg	0.8757	0.9032	0.8704	4494

Training with: lr = 0.0001, epochs = 100

```

Epoch 1/100 - Train Loss: 1.0920, Val Loss: 1.0540
Epoch 5/100 - Train Loss: 0.4577, Val Loss: 0.4395
Epoch 10/100 - Train Loss: 0.3301, Val Loss: 0.3385
Epoch 15/100 - Train Loss: 0.2950, Val Loss: 0.3079
Epoch 20/100 - Train Loss: 0.2715, Val Loss: 0.2873
Epoch 25/100 - Train Loss: 0.2511, Val Loss: 0.2691
Epoch 30/100 - Train Loss: 0.2294, Val Loss: 0.2487

```

```

Epoch 35/100 - Train Loss: 0.2082, Val Loss: 0.2281
Epoch 40/100 - Train Loss: 0.1884, Val Loss: 0.2101
Epoch 45/100 - Train Loss: 0.1742, Val Loss: 0.1970
Epoch 50/100 - Train Loss: 0.1640, Val Loss: 0.1881
Epoch 55/100 - Train Loss: 0.1565, Val Loss: 0.1816
Epoch 60/100 - Train Loss: 0.1513, Val Loss: 0.1774
Epoch 65/100 - Train Loss: 0.1472, Val Loss: 0.1729
Epoch 70/100 - Train Loss: 0.1458, Val Loss: 0.1716
Epoch 75/100 - Train Loss: 0.1422, Val Loss: 0.1687
Epoch 80/100 - Train Loss: 0.1399, Val Loss: 0.1664
Epoch 85/100 - Train Loss: 0.1388, Val Loss: 0.1653
Epoch 90/100 - Train Loss: 0.1360, Val Loss: 0.1665
Epoch 95/100 - Train Loss: 0.1347, Val Loss: 0.1612
Epoch 100/100 - Train Loss: 0.1325, Val Loss: 0.1587
Warning: deep_L5_widths_32_32_16_8_4_AdamW_0.0001_100 made no predictions for
classes: [3]

```

```

Validation report for optimizer lr_0.0001_epochs_100:
      precision    recall   f1-score   support

```

0	0.9518	0.9885	0.9698	3378
1	0.9103	0.9614	0.9352	285
2	0.9781	0.8656	0.9184	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9530	4494
macro avg	0.7101	0.7039	0.7058	4494
weighted avg	0.9416	0.9530	0.9465	4494

```

Training with: lr = 0.0001, epochs = 150
Epoch 1/150 - Train Loss: 1.3744, Val Loss: 1.3622
Epoch 5/150 - Train Loss: 1.1133, Val Loss: 1.0969
Epoch 10/150 - Train Loss: 0.9927, Val Loss: 0.9866
Epoch 15/150 - Train Loss: 0.9007, Val Loss: 0.8968
Epoch 20/150 - Train Loss: 0.8197, Val Loss: 0.8171
Epoch 25/150 - Train Loss: 0.7475, Val Loss: 0.7462
Epoch 30/150 - Train Loss: 0.6840, Val Loss: 0.6843
Epoch 35/150 - Train Loss: 0.6290, Val Loss: 0.6303
Epoch 40/150 - Train Loss: 0.5818, Val Loss: 0.5844
Epoch 45/150 - Train Loss: 0.5418, Val Loss: 0.5458
Epoch 50/150 - Train Loss: 0.5083, Val Loss: 0.5135
Epoch 55/150 - Train Loss: 0.4815, Val Loss: 0.4875
Epoch 60/150 - Train Loss: 0.4597, Val Loss: 0.4668
Epoch 65/150 - Train Loss: 0.4423, Val Loss: 0.4503
Epoch 70/150 - Train Loss: 0.4293, Val Loss: 0.4379
Epoch 75/150 - Train Loss: 0.4191, Val Loss: 0.4283
Epoch 80/150 - Train Loss: 0.4113, Val Loss: 0.4214

```

```

Epoch 85/150 - Train Loss: 0.4052, Val Loss: 0.4165
Epoch 90/150 - Train Loss: 0.4015, Val Loss: 0.4130
Epoch 95/150 - Train Loss: 0.3984, Val Loss: 0.4104
Epoch 100/150 - Train Loss: 0.3967, Val Loss: 0.4089
Epoch 105/150 - Train Loss: 0.3952, Val Loss: 0.4075
Epoch 110/150 - Train Loss: 0.3945, Val Loss: 0.4068
Epoch 115/150 - Train Loss: 0.3939, Val Loss: 0.4064
Epoch 120/150 - Train Loss: 0.3932, Val Loss: 0.4060
Epoch 125/150 - Train Loss: 0.3933, Val Loss: 0.4059
Epoch 130/150 - Train Loss: 0.3923, Val Loss: 0.4057
Epoch 135/150 - Train Loss: 0.3925, Val Loss: 0.4058
Epoch 140/150 - Train Loss: 0.3921, Val Loss: 0.4059
Epoch 145/150 - Train Loss: 0.3925, Val Loss: 0.4057
Early stopping at epoch 148 (best val loss: 0.405611)
Warning: deep_L5_widths_32_32_16_8_4_AdamW_0.0001_150 made no predictions for
classes: [1, 3]

```

```

Validation report for optimizer lr_0.0001_epochs_150:
      precision    recall   f1-score   support

```

0	0.8808	1.0000	0.9366	3378
1	0.0000	0.0000	0.0000	285
2	1.0000	0.8514	0.9197	774
3	0.0000	0.0000	0.0000	57
accuracy			0.8983	4494
macro avg	0.4702	0.4629	0.4641	4494
weighted avg	0.8343	0.8983	0.8625	4494

```

Training with: lr = 0.0001, epochs = 200
Epoch 1/200 - Train Loss: 1.3769, Val Loss: 1.3613
Epoch 5/200 - Train Loss: 1.1986, Val Loss: 1.1379
Epoch 10/200 - Train Loss: 0.9002, Val Loss: 0.8978
Epoch 15/200 - Train Loss: 0.8071, Val Loss: 0.8069
Epoch 20/200 - Train Loss: 0.5318, Val Loss: 0.4675
Epoch 25/200 - Train Loss: 0.3034, Val Loss: 0.3113
Epoch 30/200 - Train Loss: 0.2693, Val Loss: 0.2839
Epoch 35/200 - Train Loss: 0.2512, Val Loss: 0.2674
Epoch 40/200 - Train Loss: 0.2379, Val Loss: 0.2566
Epoch 45/200 - Train Loss: 0.2273, Val Loss: 0.2477
Epoch 50/200 - Train Loss: 0.2184, Val Loss: 0.2408
Epoch 55/200 - Train Loss: 0.2123, Val Loss: 0.2367
Epoch 60/200 - Train Loss: 0.2053, Val Loss: 0.2299
Epoch 65/200 - Train Loss: 0.1999, Val Loss: 0.2257
Epoch 70/200 - Train Loss: 0.1935, Val Loss: 0.2211
Epoch 75/200 - Train Loss: 0.1875, Val Loss: 0.2151
Epoch 80/200 - Train Loss: 0.1828, Val Loss: 0.2106

```

```

Epoch 85/200 - Train Loss: 0.1798, Val Loss: 0.2089
Epoch 90/200 - Train Loss: 0.1774, Val Loss: 0.2065
Epoch 95/200 - Train Loss: 0.1736, Val Loss: 0.2037
Epoch 100/200 - Train Loss: 0.1701, Val Loss: 0.2018
Epoch 105/200 - Train Loss: 0.1676, Val Loss: 0.1993
Epoch 110/200 - Train Loss: 0.1686, Val Loss: 0.1998
Epoch 115/200 - Train Loss: 0.1667, Val Loss: 0.1983
Epoch 120/200 - Train Loss: 0.1650, Val Loss: 0.1970
Epoch 125/200 - Train Loss: 0.1640, Val Loss: 0.1971
Epoch 130/200 - Train Loss: 0.1627, Val Loss: 0.1948
Epoch 135/200 - Train Loss: 0.1618, Val Loss: 0.1937
Epoch 140/200 - Train Loss: 0.1606, Val Loss: 0.1931
Epoch 145/200 - Train Loss: 0.1597, Val Loss: 0.1910
Epoch 150/200 - Train Loss: 0.1586, Val Loss: 0.1957
Epoch 155/200 - Train Loss: 0.1580, Val Loss: 0.1911
Epoch 160/200 - Train Loss: 0.1574, Val Loss: 0.1896
Epoch 165/200 - Train Loss: 0.1575, Val Loss: 0.1886
Epoch 170/200 - Train Loss: 0.1552, Val Loss: 0.1871
Epoch 175/200 - Train Loss: 0.1540, Val Loss: 0.1867
Epoch 180/200 - Train Loss: 0.1523, Val Loss: 0.1855
Epoch 185/200 - Train Loss: 0.1509, Val Loss: 0.1848
Epoch 190/200 - Train Loss: 0.1501, Val Loss: 0.1845
Epoch 195/200 - Train Loss: 0.1495, Val Loss: 0.1872
Epoch 200/200 - Train Loss: 0.1494, Val Loss: 0.1827
Warning: deep_L5_widths_32_32_16_8_4_AdamW_0.0001_200 made no predictions for
classes: [3]

```

Validation report for optimizer lr_0.0001_epochs_200:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.9422	0.9805	0.9610	3378
1	0.7973	0.8281	0.8124	285
2	0.9898	0.8734	0.9279	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9399	4494
macro avg	0.6823	0.6705	0.6753	4494
weighted avg	0.9293	0.9399	0.9337	4494

Training with: lr = 0.005, epochs = 100

```

Epoch 1/100 - Train Loss: 0.4945, Val Loss: 0.2761
Epoch 5/100 - Train Loss: 0.1516, Val Loss: 0.1746
Epoch 10/100 - Train Loss: 0.1449, Val Loss: 0.1589
Epoch 15/100 - Train Loss: 0.1279, Val Loss: 0.1336
Epoch 20/100 - Train Loss: 0.1205, Val Loss: 0.1238
Epoch 25/100 - Train Loss: 0.1222, Val Loss: 0.1280
Epoch 30/100 - Train Loss: 0.1181, Val Loss: 0.1371

```

```

Epoch 35/100 - Train Loss: 0.1137, Val Loss: 0.1135
Epoch 40/100 - Train Loss: 0.1266, Val Loss: 0.1215
Epoch 45/100 - Train Loss: 0.1192, Val Loss: 0.1420
Epoch 50/100 - Train Loss: 0.1089, Val Loss: 0.1266
Early stopping at epoch 55 (best val loss: 0.113513)
Warning: deep_L5_widths_32_32_16_8_4_AdamW_0.005_100 made no predictions for
classes: [3]

```

Validation report for optimizer lr_0.005_epochs_100:				
	precision	recall	f1-score	support
0	0.9572	0.9858	0.9713	3378
1	0.8799	0.9509	0.9140	285
2	0.9830	0.8979	0.9386	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9559	4494
macro avg	0.7050	0.7087	0.7060	4494
weighted avg	0.9446	0.9559	0.9497	4494

Training with: lr = 0.005, epochs = 150

```

Epoch 1/150 - Train Loss: 0.4697, Val Loss: 0.2820
Epoch 5/150 - Train Loss: 0.1766, Val Loss: 0.1966
Epoch 10/150 - Train Loss: 0.1599, Val Loss: 0.1808
Epoch 15/150 - Train Loss: 0.1549, Val Loss: 0.1663
Epoch 20/150 - Train Loss: 0.1528, Val Loss: 0.1688
Epoch 25/150 - Train Loss: 0.1472, Val Loss: 0.1661
Epoch 30/150 - Train Loss: 0.1430, Val Loss: 0.1602
Epoch 35/150 - Train Loss: 0.1380, Val Loss: 0.1639
Epoch 40/150 - Train Loss: 0.1441, Val Loss: 0.1535
Epoch 45/150 - Train Loss: 0.1398, Val Loss: 0.1576
Epoch 50/150 - Train Loss: 0.1392, Val Loss: 0.2838
Epoch 55/150 - Train Loss: 0.1382, Val Loss: 0.1553
Epoch 60/150 - Train Loss: 0.1392, Val Loss: 0.1437
Epoch 65/150 - Train Loss: 0.1359, Val Loss: 0.1536
Epoch 70/150 - Train Loss: 0.1371, Val Loss: 0.1484
Epoch 75/150 - Train Loss: 0.1301, Val Loss: 0.1439
Epoch 80/150 - Train Loss: 0.1264, Val Loss: 0.1636
Epoch 85/150 - Train Loss: 0.1290, Val Loss: 0.1598
Epoch 90/150 - Train Loss: 0.1287, Val Loss: 0.1378
Epoch 95/150 - Train Loss: 0.1331, Val Loss: 0.1515
Epoch 100/150 - Train Loss: 0.1254, Val Loss: 0.1480
Early stopping at epoch 103 (best val loss: 0.135626)

```

Validation report for optimizer lr_0.005_epochs_150:				
	precision	recall	f1-score	support

0	0.9506	0.9864	0.9682	3378
1	0.9247	0.9474	0.9359	285
2	0.9985	0.8514	0.9191	774
3	0.3514	0.2281	0.2766	57
accuracy			0.9510	4494
macro avg	0.8063	0.7533	0.7749	4494
weighted avg	0.9496	0.9510	0.9489	4494

Training with: lr = 0.005, epochs = 200
Epoch 1/200 - Train Loss: 0.4412, Val Loss: 0.2837
Epoch 5/200 - Train Loss: 0.1862, Val Loss: 0.2007
Epoch 10/200 - Train Loss: 0.1598, Val Loss: 0.1717
Epoch 15/200 - Train Loss: 0.1406, Val Loss: 0.1455
Epoch 20/200 - Train Loss: 0.1329, Val Loss: 0.1495
Epoch 25/200 - Train Loss: 0.1258, Val Loss: 0.1322
Epoch 30/200 - Train Loss: 0.1194, Val Loss: 0.1463
Epoch 35/200 - Train Loss: 0.1180, Val Loss: 0.1220
Epoch 40/200 - Train Loss: 0.1215, Val Loss: 0.1406
Epoch 45/200 - Train Loss: 0.1151, Val Loss: 0.1390
Epoch 50/200 - Train Loss: 0.1181, Val Loss: 0.1251
Early stopping at epoch 53 (best val loss: 0.118734)

Validation report for optimizer lr_0.005_epochs_200:				
	precision	recall	f1-score	support
0	0.9658	0.9870	0.9763	3378
1	0.9343	0.9474	0.9408	285
2	0.9884	0.8811	0.9317	774
3	0.5397	0.5965	0.5667	57
accuracy			0.9613	4494
macro avg	0.8570	0.8530	0.8539	4494
weighted avg	0.9623	0.9613	0.9612	4494

Q: Now, focus on the architecture with the best optimizer. Evaluate the effects of the different learning rates and epochs. Report the test results for the best model.

```
[112]: best_lr_epochs_model_tag = 'lr_0.005_epochs_200'
model = trained_lr_epochs_models[best_lr_epochs_model_tag]

print(f"\nTest set classification report ({best_lr_epochs_model_tag}):")
report = evaluate_model(model, X_test_tensor_no_port, y_test_no_port)
print(report)
```

Test set classification report (lr_0.005_epochs_200):

	precision	recall	f1-score	support
0	0.9664	0.9873	0.9767	3378
1	0.9281	0.9476	0.9377	286
2	0.9927	0.8797	0.9328	773
3	0.5000	0.5789	0.5366	57
accuracy			0.9611	4494
macro avg	0.8468	0.8484	0.8460	4494
weighted avg	0.9626	0.9611	0.9611	4494

1.7 Task 6 — Overfitting and Regularization

We analyze overfitting and apply regularization techniques to improve generalization.

Base model: - Layers: **6** - Hidden widths: [256, 128, 64, 32, 16] - Activation: **ReLU** - Optimizer: **AdamW**, learning rate = 5e-4 - Batch size: **128** - Epochs: up to 50

Steps: - Train baseline model and inspect loss curves for overfitting. - Add **Dropout**, **Batch Normalization**, and **Weight Decay** (AdamW regularization). - Compare validation/test performance across variants. - Discuss which regularization technique best mitigates overfitting.

```
[113]: # Create directory for plots
save_dir = results_path + 'images/' + 'task6_plots/'
os.makedirs(save_dir, exist_ok=True)
```

1.7.1 Training

```
[114]: # --- Define 6-layer FFNN with Regularization Options ---
```

```
class RegularizedDeepFFNN(nn.Module):
    def __init__(self, input_dim, layer_widths, output_dim, activation='relu', ↴
                 dropout_prob=0.0, use_batchnorm=False):
        super(RegularizedDeepFFNN, self).__init__()
        layers = []
        prev_width = input_dim
        for i, width in enumerate(layer_widths):
            layers.append(nn.Linear(prev_width, width))
            if use_batchnorm:
                layers.append(nn.BatchNorm1d(width))
            if activation == 'relu':
                layers.append(nn.ReLU())
            if dropout_prob > 0 and i < len(layer_widths) - 1: # Apply dropout ↴
                to hidden layers
                layers.append(nn.Dropout(dropout_prob))
            prev_width = width
        layers.append(nn.Linear(prev_width, output_dim))
        self.net = nn.Sequential(*layers)
```

```

def forward(self, x):
    return self.net(x)

```

```

[115]: # --- Experiment with Regularization ---

input_dim_reg = X_train_std_no_port.shape[1]
output_dim_reg = len(np.unique(y_train_no_port))
layer_widths_reg = [256, 128, 64, 32, 16] # Same widths as the baseline 6-layer
    ↵model

# Define hyperparameters for regularization experiments
base_lr = 0.0005
epochs_reg = 50
batch_size_reg = 128
min_delta_reg = 0.00001
patience_reg = 20

regularization_configs = {
    'Baseline': # No dropout, no batch norm, no weight decay
        {'dropout_prob': 0.0, 'use_batchnorm': False, 'weight_decay': 0.0},
    'Dropout_0.5': # Dropout 50% (no BN, no weight decay)
        {'dropout_prob': 0.5, 'use_batchnorm': False, 'weight_decay': 0.0},
    'BatchNorm': # Batch normalization only
        {'use_batchnorm': True, 'dropout_prob': 0.0, 'weight_decay': 0.0},
    'BatchNorm_Dropout_0.5': # BatchNorm + Dropout 50%
        {'use_batchnorm': True, 'dropout_prob': 0.5, 'weight_decay': 0.0},
    'WeightDecay_1e-4': # L2 regularization via weight_decay=1e-4 (no dropout,
    ↵no BN)
        {'dropout_prob': 0.0, 'use_batchnorm': False, 'weight_decay': 1e-4},
    'WeightDecay_1e-4_BN_Dropout_0.5': # All three: BatchNorm + Dropout 50% +
    ↵Weight Decay 1e-4
        {'use_batchnorm': True, 'dropout_prob': 0.5, 'weight_decay': 1e-4}
}

regularized_models = {}
regularized_loss_curves = {}
regularized_validation_reports = {}
regularized_test_reports = {}

for config_name, params in regularization_configs.items():
    print(f"\nTraining model with {config_name}...")

    # Instantiate model with regularization
    model_reg = RegularizedDeepFFNN(
        input_dim_reg,
        layer_widths_reg,

```

```

        output_dim_reg,
        activation='relu',
        dropout_prob=params.get('dropout_prob', 0.0),
        use_batchnorm=params.get('use_batchnorm', False)
    )

# Define optimizer with weight decay
optimizer_reg = optim.AdamW(model_reg.parameters(), lr=base_lr, weight_decay=params.get('weight_decay', 0.0))
criterion = nn.CrossEntropyLoss()

# Create DataLoaders (using the same as the baseline 6-layer model)
train_loader_reg = train_loader_no_port
val_loader_reg = val_loader_no_port

# Move model to device
model_reg = model_reg.to(device)

# Training
model_reg, train_loss_reg, val_loss_reg = train_model(
    model_reg,
    train_loader_reg,
    val_loader_reg,
    epochs_reg,
    optimizer_reg,
    criterion,
    min_delta=min_delta_reg,
    patience=patience_reg
)

regularized_models[config_name] = model_reg
regularized_loss_curves[config_name] = (train_loss_reg, val_loss_reg)

model_name = f"deep_L3_widths_{'_'.join(map(str, layer_widths_reg))}_reg_{config_name}"

# Evaluate on validation set
report_val_reg = evaluate_model(model_reg, X_val_tensor_no_port, y_val_no_port, model_name)
regularized_validation_reports[config_name] = report_val_reg
print(f"\nValidation report for {config_name}:")
print(report_val_reg)

# Evaluate on test set
report_test_reg = evaluate_model(model_reg, X_test_tensor_no_port, y_test_no_port)
regularized_test_reports[config_name] = report_test_reg

```

```

print(f"\nTest report for {config_name}:")
print(report_test_reg)

# It is now possible compare the reports in regularized_validation_reports and
regularized_test_reports

```

Training model with Baseline...

```

Epoch 1/50 - Train Loss: 0.5558, Val Loss: 0.3166
Epoch 5/50 - Train Loss: 0.1646, Val Loss: 0.1904
Epoch 10/50 - Train Loss: 0.1322, Val Loss: 0.1489
Epoch 15/50 - Train Loss: 0.1262, Val Loss: 0.1447
Epoch 20/50 - Train Loss: 0.1172, Val Loss: 0.1387
Epoch 25/50 - Train Loss: 0.1083, Val Loss: 0.1213
Epoch 30/50 - Train Loss: 0.1034, Val Loss: 0.1297
Epoch 35/50 - Train Loss: 0.1006, Val Loss: 0.1168
Epoch 40/50 - Train Loss: 0.1018, Val Loss: 0.1274
Epoch 45/50 - Train Loss: 0.1023, Val Loss: 0.1293
Epoch 50/50 - Train Loss: 0.1022, Val Loss: 0.1188

```

Validation report for Baseline:

	precision	recall	f1-score	support
0	0.9669	0.9843	0.9755	3378
1	0.9310	0.9474	0.9391	285
2	0.9789	0.8992	0.9374	774
3	0.6296	0.5965	0.6126	57
accuracy			0.9624	4494
macro avg	0.8766	0.8568	0.8662	4494
weighted avg	0.9624	0.9624	0.9620	4494

Test report for Baseline:

	precision	recall	f1-score	support
0	0.9667	0.9876	0.9770	3378
1	0.9249	0.9476	0.9361	286
2	0.9817	0.9017	0.9400	773
3	0.7750	0.5439	0.6392	57
accuracy			0.9646	4494
macro avg	0.9121	0.8452	0.8731	4494
weighted avg	0.9642	0.9646	0.9638	4494

Training model with Dropout_0.5...

```

Epoch 1/50 - Train Loss: 0.6715, Val Loss: 0.3437

```

Epoch 5/50 - Train Loss: 0.2813, Val Loss: 0.2687
Epoch 10/50 - Train Loss: 0.2255, Val Loss: 0.2099
Epoch 15/50 - Train Loss: 0.1859, Val Loss: 0.1757
Epoch 20/50 - Train Loss: 0.1659, Val Loss: 0.1581
Epoch 25/50 - Train Loss: 0.1606, Val Loss: 0.1517
Epoch 30/50 - Train Loss: 0.1526, Val Loss: 0.1444
Epoch 35/50 - Train Loss: 0.1478, Val Loss: 0.1402
Epoch 40/50 - Train Loss: 0.1499, Val Loss: 0.1389
Epoch 45/50 - Train Loss: 0.1496, Val Loss: 0.1419
Epoch 50/50 - Train Loss: 0.1462, Val Loss: 0.1339
Warning: deep_L3_widths_256_128_64_32_16_reg_Dropout_0.5 made no predictions for classes: [3]

Validation report for Dropout_0.5:

	precision	recall	f1-score	support
0	0.9474	0.9802	0.9635	3378
1	0.8108	0.9474	0.8738	285
2	0.9925	0.8540	0.9181	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9439	4494
macro avg	0.6877	0.6954	0.6888	4494
weighted avg	0.9345	0.9439	0.9377	4494

Warning: Unnamed model made no predictions for classes: [3]

Test report for Dropout_0.5:

	precision	recall	f1-score	support
0	0.9440	0.9772	0.9603	3378
1	0.7890	0.9545	0.8639	286
2	0.9939	0.8370	0.9087	773
3	0.0000	0.0000	0.0000	57
accuracy			0.9393	4494
macro avg	0.6817	0.6922	0.6832	4494
weighted avg	0.9307	0.9393	0.9331	4494

Training model with BatchNorm...

Epoch 1/50 - Train Loss: 0.5386, Val Loss: 0.3429
Epoch 5/50 - Train Loss: 0.1481, Val Loss: 0.1696
Epoch 10/50 - Train Loss: 0.1550, Val Loss: 0.1514
Epoch 15/50 - Train Loss: 0.1409, Val Loss: 0.1518
Epoch 20/50 - Train Loss: 0.1426, Val Loss: 0.1577
Epoch 25/50 - Train Loss: 0.1333, Val Loss: 0.1451
Epoch 30/50 - Train Loss: 0.1336, Val Loss: 0.1594

Epoch 35/50 - Train Loss: 0.1299, Val Loss: 0.1726
Epoch 40/50 - Train Loss: 0.1284, Val Loss: 0.1735
Epoch 45/50 - Train Loss: 0.1320, Val Loss: 0.1424
Epoch 50/50 - Train Loss: 0.1269, Val Loss: 0.2123

Validation report for BatchNorm:

	precision	recall	f1-score	support
0	0.9585	0.9855	0.9718	3378
1	0.9278	0.9474	0.9375	285
2	0.9926	0.8721	0.9285	774
3	0.3600	0.3158	0.3364	57
accuracy			0.9551	4494
macro avg	0.8098	0.7802	0.7936	4494
weighted avg	0.9549	0.9551	0.9541	4494

Test report for BatchNorm:

	precision	recall	f1-score	support
0	0.9600	0.9870	0.9733	3378
1	0.9128	0.9510	0.9315	286
2	0.9941	0.8719	0.9290	773
3	0.4889	0.3860	0.4314	57
accuracy			0.9573	4494
macro avg	0.8389	0.7990	0.8163	4494
weighted avg	0.9569	0.9573	0.9561	4494

Training model with BatchNorm_Dropout_0.5...

Epoch 1/50 - Train Loss: 0.8901, Val Loss: 0.5737
Epoch 5/50 - Train Loss: 0.2895, Val Loss: 0.2536
Epoch 10/50 - Train Loss: 0.2453, Val Loss: 0.2137
Epoch 15/50 - Train Loss: 0.2246, Val Loss: 0.1937
Epoch 20/50 - Train Loss: 0.2197, Val Loss: 0.1932
Epoch 25/50 - Train Loss: 0.2100, Val Loss: 0.1848
Epoch 30/50 - Train Loss: 0.2054, Val Loss: 0.1796
Epoch 35/50 - Train Loss: 0.1948, Val Loss: 0.1747
Epoch 40/50 - Train Loss: 0.1940, Val Loss: 0.1717
Epoch 45/50 - Train Loss: 0.1935, Val Loss: 0.1707
Epoch 50/50 - Train Loss: 0.1922, Val Loss: 0.1737
Warning: deep_L3_widths_256_128_64_32_16_reg_BatchNorm_Dropout_0.5 made no predictions for classes: [3]

Validation report for BatchNorm_Dropout_0.5:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.9464	0.9725	0.9593	3378
1	0.7459	0.9474	0.8346	285
2	0.9970	0.8514	0.9185	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9377	4494
macro avg	0.6723	0.6928	0.6781	4494
weighted avg	0.9304	0.9377	0.9322	4494

Warning: Unnamed model made no predictions for classes: [3]

Test report for BatchNorm_Dropout_0.5:

	precision	recall	f1-score	support
0	0.9431	0.9710	0.9568	3378
1	0.7378	0.9545	0.8323	286
2	0.9985	0.8344	0.9091	773
3	0.0000	0.0000	0.0000	57
accuracy			0.9341	4494
macro avg	0.6698	0.6900	0.6746	4494
weighted avg	0.9276	0.9341	0.9286	4494

Training model with WeightDecay_1e-4...

Epoch 1/50 - Train Loss: 0.5557, Val Loss: 0.3046
 Epoch 5/50 - Train Loss: 0.1690, Val Loss: 0.1760
 Epoch 10/50 - Train Loss: 0.1381, Val Loss: 0.1736
 Epoch 15/50 - Train Loss: 0.1263, Val Loss: 0.1565
 Epoch 20/50 - Train Loss: 0.1294, Val Loss: 0.1530
 Epoch 25/50 - Train Loss: 0.1265, Val Loss: 0.1441
 Epoch 30/50 - Train Loss: 0.1212, Val Loss: 0.1385
 Epoch 35/50 - Train Loss: 0.1111, Val Loss: 0.1243
 Epoch 40/50 - Train Loss: 0.1143, Val Loss: 0.1333
 Epoch 45/50 - Train Loss: 0.1119, Val Loss: 0.1220
 Epoch 50/50 - Train Loss: 0.1073, Val Loss: 0.1305

Validation report for WeightDecay_1e-4:

	precision	recall	f1-score	support
0	0.9695	0.9787	0.9741	3378
1	0.9340	0.9439	0.9389	285
2	0.9588	0.9031	0.9301	774
3	0.3881	0.4561	0.4194	57
accuracy			0.9568	4494
macro avg	0.8126	0.8204	0.8156	4494

weighted avg	0.9580	0.9568	0.9572	4494
--------------	--------	--------	--------	------

Test report for WeightDecay_1e-4:

	precision	recall	f1-score	support
0	0.9707	0.9796	0.9751	3378
1	0.9218	0.9476	0.9345	286
2	0.9652	0.8978	0.9303	773
3	0.4722	0.5965	0.5271	57
accuracy			0.9586	4494
macro avg	0.8325	0.8554	0.8418	4494
weighted avg	0.9603	0.9586	0.9591	4494

Training model with WeightDecay_1e-4_BN_Dropout_0.5...

Epoch 1/50 - Train Loss: 0.9748, Val Loss: 0.6319

Epoch 5/50 - Train Loss: 0.3131, Val Loss: 0.2765

Epoch 10/50 - Train Loss: 0.2544, Val Loss: 0.2077

Epoch 15/50 - Train Loss: 0.2302, Val Loss: 0.1894

Epoch 20/50 - Train Loss: 0.2205, Val Loss: 0.1799

Epoch 25/50 - Train Loss: 0.2161, Val Loss: 0.1779

Epoch 30/50 - Train Loss: 0.2024, Val Loss: 0.1749

Epoch 35/50 - Train Loss: 0.1964, Val Loss: 0.1680

Epoch 40/50 - Train Loss: 0.1988, Val Loss: 0.1711

Epoch 45/50 - Train Loss: 0.1939, Val Loss: 0.1716

Epoch 50/50 - Train Loss: 0.1987, Val Loss: 0.1708

Warning: deep_L3_widths_256_128_64_32_16_reg_WeightDecay_1e-4_BN_Dropout_0.5

made no predictions for classes: [3]

Validation report for WeightDecay_1e-4_BN_Dropout_0.5:

	precision	recall	f1-score	support
0	0.9455	0.9766	0.9608	3378
1	0.7733	0.9333	0.8458	285
2	0.9970	0.8514	0.9185	774
3	0.0000	0.0000	0.0000	57
accuracy			0.9399	4494
macro avg	0.6789	0.6903	0.6813	4494
weighted avg	0.9315	0.9399	0.9340	4494

Warning: Unnamed model made no predictions for classes: [3]

Test report for WeightDecay_1e-4_BN_Dropout_0.5:

precision recall f1-score support

0	0.9434	0.9725	0.9577	3378
1	0.7479	0.9545	0.8387	286
2	0.9985	0.8357	0.9099	773
3	0.0000	0.0000	0.0000	57
accuracy			0.9355	4494
macro avg	0.6725	0.6907	0.6766	4494
weighted avg	0.9285	0.9355	0.9298	4494

1.7.2 Evaluating


[116]: # --- Plot loss curves for all regularized models ---

```
for config_name, (train_loss, val_loss) in regularized_loss_curves.items():
    plt.figure(figsize=(10,6))
    plt.plot(train_loss, label=f'Train Loss - {config_name}')
    plt.plot(val_loss, label=f'Val Loss - {config_name}')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.title(f'Loss Curves - {config_name}')
    plt.legend()


    # Save the plot to the specified path
    save_plot(plt.gcf(), f"{config_name}_loss_curve", save_dir) # Use plt.gcf() ↴
    ↴ to get the current figure

plt.show()
```


Saved plot: ../results/images/task6_plots/Baseline_loss_curve.png




Saved plot: ./results/images/task6_plots/Dropout_0.5_loss_curve.png




Saved plot: ../results/images/task6_plots/BatchNorm_loss_curve.png




Saved plot: ../results/images/task6_plots/BatchNorm_Dropout_0.5_loss_curve.png



Saved plot: ../results/images/task6_plots/WeightDecay_1e-4_loss_curve.png



Saved plot:
../results/images/task6_plots/WeightDecay_1e-4_BN_Dropout_0.5_loss_curve.png



Q: Each Neurons per Layer describes how many neurons must be present in each layer. For example: 1st layer 256, 2nd layer 128 etc. What do the losses look like? Is the model overfitting? The training and validation losses both decrease smoothly and stabilize around 0.10-0.12, showing consistent convergence. The validation loss stays slightly higher than the training loss, which is expected — this indicates good generalization rather than overfitting.

If the model were overfitting, we would see the training loss continue to drop while the validation loss increased or fluctuated strongly. Here, both curves follow the same trend and plateau together.

The losses show a healthy training process — the model converges well and does not overfit, maintaining high validation accuracy (~96%).

Q: Now apply normalization techniques (dropout, batch normalization) and play with the regularization of the weights (AdamW's weight decay). What impact do the different normalization techniques have on validation and testing performance? Impact summary:

- **Baseline (AdamW)** - Best overall (Val 96%, Test 96%). Smooth convergence, no overfitting, and correct predictions for all classes.
- **Dropout (0.5)** - Validation loss a bit lower than training; minority class (3) never predicted.
- **BatchNorm** - Validation loss unstable, signs of overfitting/instability; poor minority class recall.

- **BatchNorm + Dropout (0.5)** - Too much regularization → strong underfitting; class 3 ignored.
- **Weight Decay (1e-4)** - Slightly improves generalization and keeps losses stable; close to baseline.
- **Weight Decay + BN + Dropout (0.5)** - Over-regularized; underfits, poor validation.

Best setup: *AdamW + small weight decay (1e-4)*.

Too strong: *Dropout / BatchNorm* on this tabular dataset → underfitting, unstable validation.
